Scrap Your Boilerplate with Object Algebras

Haoyuan Zhang   Zewei Chu
Bruno C. d. S. Oliveira

The University of Hong Kong, China
{hyzhang,jason91,bruno}@cs.hku.hk

Tijs van der Storm

Centrum Wiskunde & Informatica (CWI), The
Netherlands
storm@cwi.nl

Abstract

Traversing complex *Abstract Syntax Trees* (ASTs) typically requires large amounts of tedious boilerplate code. For many operations most of the code simply walks the structure, and only a small portion of the code implements the functionality that motivated the traversal in the first place. This paper presents a type-safe Java framework called **Shy** that removes much of this boilerplate code. In **Shy** *Object Algebras* are used to describe complex and extensible AST structures. Using Java annotations **Shy** generates generic boilerplate code for various types of traversals. For a concrete traversal, users of **Shy** can then inherit from the generated code and override only the interesting cases. Consequently, the amount of code that users need to write is significantly smaller. Moreover, traversals using the **Shy** framework are also much more *structure shy*, becoming more adaptive to future changes or extensions to the AST structure. To prove the effectiveness of the approach, we applied **Shy** in the implementation of a domain-specific questionnaire language. Our results show that for a large number of traversals there was a significant reduction in the amount of user-defined code.

1. Introduction

Various language processing tools or libraries for programming languages, domain-specific languages, mark-up languages like HTML, or data-interchange languages like XML or JSON require complex *Abstract Syntax Tree* (AST) structures. In those applications ASTs are the key data structure needed to model the various constructs of these languages. Such ASTs have various different types of nodes, which can range from a few dozen to several hundred kinds of nodes (for example in the ASTs of languages like Java or Cobol).

Static types are helpful to deal with such complex ASTs. Static types formalize the distinction between different kinds of nodes. Furthermore the distinctions are helpful to ensure that traversals over these ASTs have an appropriate piece of code that deals with each different type of node. This can prevent a large class of run-time errors that would not be detected otherwise.

Unfortunately, when traversing such ASTs, the number of nodes and the enforced type distinctions between nodes can lead to so-called *boilerplate code* [11]: code that is similar for most types of nodes and which essentially just walks the structure. Operations where the "interesting" code is limited to only a small portion of nodes are called *structure shy* [17]. A typical example is computing the free variables of an expression in some programming language. In this case, the interesting code occurs in the nodes representing the binding and variable constructs. In all other cases, the code would just deal with walking the structure. In ASTs with dozens or hundreds of kinds of nodes, having to explicitly write code for each kind of node is both tedious and error-prone.

The boilerplate problem in implementing traversals has received considerable attention in the past. For example, both *Adaptive Object-Oriented Programming* (AOOP) [17] and *Strategic Programming* [3, 26] are aimed partly at solving this problem. Most approaches to AOOP and strategic programming use some meta-programming techniques, such as code generation or reflection. The use of meta-programming offers programmers an easy way to avoid having to write boilerplate code. This has important benefits: users have to write much less code; and the code becomes much more adaptive to changes. Nevertheless, such meta-

programming based approaches usually come at the cost of other desirable properties, such as type-safety, extensibility or separate compilation. The functional programming community has also studied the problem. For instance, the popular "*Scrap your boilerplate*" [11] approach supports type-safety and separate compilation. Most of the techniques used in functional languages, however, cannot be easily ported to mainstream OO languages like Java, and are limited in terms of extensibility.

This paper presents a Java framework called **Shy** that allows users to define type-safe and extensible structure-shy operations. **Shy** uses *Object Algebras* [19] to describe ASTs, and to write operations over ASTs in a style similar to writing folds in functional programming. However unlike standard folds in functional programming, Object Algebras are *extensible*: when new kinds of nodes are introduced in the data type, the operations can be extended as well, without changing existing code.

In **Shy** Object Algebra interfaces are combined with Java annotations to generate generic traversal code automatically. The generated code accounts for four different types of traversals: *queries*; *transformations*; *generalized queries*; and *contextual transformations*. Each of these four types of traversals is implemented as an Object Algebra. Programmers who want to implement structure-shy traversals can then inherit from one of these four Object Algebras, and override only the cases that deal with the interesting parts of the traversal. Consequently traversals written in **Shy** are:

- **Small in size:** With **Shy** the amount of code that programmers need to write a structure-shy traversal is significantly smaller. Often traversals in **Shy** are implementable in just a few lines of code, even for complex ASTs with hundreds of different types of nodes.
- **Adaptive and structure shy:** Traversals written with **Shy** can omit boilerplate code, making traversals more adaptive to future changes or extensions to the data type.
- **Type-safe: Shy** traversals are directly written in Java and the Java type-system ensures type-safety. No run-time casts are needed for generic traversal code or for user-defined traversal code.
- **Extensible with separate compilation:** Traversals inherit type-safe extensibility from Object Algebras. Both traversals and the AST structures are extensible. Thus it is possible to reuse traversal code for ASTs extended with additional node types. Furthermore **Shy** traversals support separate compilation.
- **Implemented in plain Java: Shy** traversals do not require a new tool or language. The approach is library based and only uses Java annotations.

To prove the effectiveness of the approach, we have applied **Shy** in the implementation of QL, a domain-specific language (DSL) for defining questionnaires that has been used before as a benchmark language [7, 9]. Our results (see de-tails in Section 11) show that for a large number of traversals there was a significant reduction in the amount of user-defined code: only 4% to 21% of the AST cases had to be implemented in comparison with code written without **Shy**.

Although **Shy**'s functional programming inspired idioms are new to mainstream Java programmers, the use of standard Java constructs and techniques makes **Shy** easy to implement, integrate in existing environments, and use. Moreover, although Java was chosen as the implementation language for **Shy**, our approach should apply to any object-oriented language with support for generics and annotations.

In summary, the contributions of this paper are:

- **Design patterns for generic traversals.** We provide a set of design patterns for various types of traversals using Object Algebras, including: *queries*, *transformations*, *generalized queries* and *contextual transformations*.
- **The Shy Java framework.** We show that those design patterns can be automatically generated for a given Object Algebra interface. The **Shy** framework[1] realizes this idea by using Java annotations to automatically generate generic traversals.
- **Applications and examples:** We show various examples of traversals, such as free variables and substitution operations, implemented concisely with **Shy**. Moreover we illustrate how certain kinds of *desugarings* can be implemented using transformations, and how multiple transformations can be chained together in a *pipeline*.
- **Case study.** We illustrate the benefits of **Shy** using a case study based on the QL domain-specific language. The results of our case study show significant savings in terms of user-defined traversal code. The case study also shows that **Shy** does not incur significant performance overhead compared to a regular AST-based implementation.

2. Background: Object Algebras

Object Algebras capture a design pattern to solve the Expression Problem [29]. As a result Object Algebras support modular and type-safe extensibility in two dimensions: data variants and operations. Object Algebras can be used conveniently to model recursive data types, such as ASTs. Here we briefly introduce Object Algebras using the example of a simple expression language.

```java
interface ExpAlg<E> {
  E Lit(int n);
  E Add(E l, E r);
}
```

The above code shows an *Object Algebra interface* which models literal and addition expressions. The interface resembles an abstract factory interface, but instead of constructing objects of a concrete type, the result is abstract through the type parameter E. A concrete Object Algebra will implement such an interface by instantiating the type parameter with

---

[1] **Available at:** https://github.com/JasonCHU/SYBwithOA

a type that constructs objects for a particular purpose. For instance, evaluation of expressions can be implemented as follows:

```java
class Eval implements ExpAlg<Integer> {
 public Integer Lit(int n) {
  return n;
 }
 public Integer Add(Integer l, Integer r) {
  return l + r;
 }
}
```

Using this algebra to create expressions will immediately evaluate the expression to its result. To add another operation the interface `ExpAlg` is implemented again. As an example, the class `Print` presented below shows a simple printing operation on expressions. Both examples illustrate extension in the dimension of operations.

```java
class Print implements ExpAlg<String> {
 public String Lit(int n) { return "" + n; }
 public String Add(String l, String r) {
  return l + " + " + r;
 }
}
```

The other dimension of extensibility (adding data variants) is realized by extending the Object Algebra interface itself. For instance, the interface `MulAlg` defined below extends the original `ExpAlg` with multiplication expressions.

```java
interface MulAlg<E> extends ExpAlg<E> {
 E Mul(E l, E r);
}
```

The key feature of Object Algebras is that the existing operations (evaluation and printing) can now be extended to support multiplication without having to change the existing `Eval` and `Print` classes. For example, evaluation is extended with multiplication as follows:

```java
class MulEval extends Eval
 implements MulAlg<Integer> {
 public Integer Mul(Integer l, Integer r) {
  return l * r;
 }
}
```

Client code The client code below shows how structures are created and operations are applied using Object Algebras. For example, The generic interface `ExpAlg` is used to create a structure as follows:

```java
<E> E makeExp(ExpAlg<E> alg) {
 return alg.Add(alg.Lit(2), alg.Lit(3));
}
```

Such a structure represents the expression "2 + 3". Similarly, the extended interface `MulAlg` supporting multiplication creates the expression "(2 + 3) * 4" as follows:

```java
<E> E makeMul(MulAlg<E> alg) {
 return alg.Mul(alg.Add(alg.Lit(2), alg.Lit(3)),
    alg.Lit(4));
}
```

```
form DriverLicense {
  name: "What is your name?" string
  age: "What is your age?" integer
  if (age >= 18)
    license: "Have a driver's license?" boolean
}
```

Figure 1. Example QL questionnaire: driver's license.

Users can then pass instances of Object Algebras to those structures for applying operations:

```java
println(makeExp(new Eval()));
println(makeExp(new Print()));
println(makeMul(new MulEval()));
```

The results are "5", "2 + 3" and "20", respectively.

In summary, Object Algebra interfaces describe recursive data types, and implementations of those interfaces represent operations. Extending the interface allows developers to extend the data type with new data variants, and (re)implementing the interface allows developers to define operations. Finally, extension in both dimensions is fully type-safe and does not compromise separate compilation.

3. An Overview of Shy

This section starts by illustrating the problem of boilerplate code when implementing traversals of complex structures. It then shows how **Shy** addresses the problem using a combination of Object Algebras [19] and Java annotations.

3.1 Traversing Object-Oriented ASTs

We start by introducing the problem of boilerplate code by considering a simplified variant of the QL language used in our case study [9], called MiniQL. Just like QL, MiniQL can be used to describe interactive questionnaires.

An example MiniQL program is shown in Figure 1. The questionnaire first asks for the user's name and age, and then, if the age is greater than or equal to 18, asks if the user has a driver's license. Because of the conditional construct, the last question will only appear when the user is actually eligible to have driver's license.

MiniQL's abstract syntax contains forms, statements (if-then and question) and expressions (only literals, variables and greater-than-or-equal). A traditional OO implementation is shown on the left Figure 2. A form (class `Form`) has a name and consists of a list of statements. Statements are conditionals (`If`) which contain an expression and a statement body, and questions (`Question`) which have a name, label and type. Expressions are standard, but limited to literals (`Lit`), variables (`Var`) and greater-than-or-equal (`GEq`).

The code in Figure 2 also shows a query over MiniQL structures, namely the collection of used variables. The operation is defined using the method `usedVars`, declared in the abstract superclasses `Stmt` and `Exp` (omitted for brevity), and implemented in the concrete statement and expression classes. As can be seen, the only interesting bit of code is

```
class Form {
  String name; List<Stmt> body;
  Set<String> usedVars() {
    Set<String> vars = new HashSet<>();
    body.forEach(s -> vars.addAll(s.usedVars()));
    return vars;
  }
}

class If extends Stmt {
  Exp cond; Stmt then;
  Set<String> usedVars() {
    Set<String> vars=new HashSet<>(cond.usedVars());
    vars.addAll(then.usedVars());
    return vars;
  }
}

class Question extends Stmt {
  String name, label, type;
  Set<String> usedVars() { return emptySet(); }
}

class Lit extends Exp {
  int n;
  Set<String> usedVars() { return emptySet(); }
}

class Var extends Exp {
  String x;
  Set<String> usedVars() { return singleton(x); }
}

class GEq extends Exp {
  Exp lhs, rhs;
  Set<String> usedVars() {
    Set<String> vars = new HashSet<>(lhs.usedVars());
    vars.addAll(rhs.usedVars());
    return vars;
  }
}
```

```
class UsedVars implements
    QLAlg<Set<String>, Set<String>, Set<String>> {

  Set<String> Form(String n, List<Set<String>> b) {
    Set<String> vars = new HashSet<>();
    b.forEach(s -> vars.addAll(s));
    return vars;
  }

  Set<String> If(Set<String> c, Set<String> t) {
    Set<String> vars = new HashSet<>(c);
    vars.addAll(t);
    return vars;
  }

  Set<String> Question(String n,String l,String t) {
    return Collections.emptySet();
  }

  Set<String> Lit(int x) {
    return Collections.emptySet();
  }

  Set<String> Var(String x) {
    return Collections.singleton(x);
  }

  Set<String> GEq(Set<String> l, Set<String> r) {
    Set<String> vars = new HashSet<>(l);
    vars.addAll(r);
    return vars;
  }
}
```

Figure 2. Implementing the "used variables" operation using traditional ASTs (left) and Object Algebras (right).

the usedVars method in class Var. All other implementations merely deal with aggregating results of their child nodes, or returning a default empty set.

The boilerplate code exhibited in the usedVars query often also applies to transformations. Consider for example a rename transformation which takes a Form and returns another form where the occurrences of the specified variable are renamed. Again, the only interesting cases would be in the Var and Question classes, where the actual renaming is applied. All other classes, however, require boilerplate to recreate the structure. The full code in Appendix A.1.1 contains a simple example of such a rename operation as well.

In addition to the significant amount of boilerplate code, there is another drawback to the traditional OO solution, which is that it does not support extensibility along the di-

mension of operations. Each new operation requires pervasive changes across the AST classes.

3.2 Modeling MiniQL with Object Algebras

The right-hand side of Figure 2 shows the used variables operation implemented using Object Algebras. The operation is a class implementing the Object Algebra interface (QLAlg) shown in Figure 3.

The UsedVars class provides an implementation for each of the methods in the object algebra interface, which together define the full used variables operation. Since the result of collecting those variables is Set<String>, all the type parameters are set to that type. Most of the method implementations simply traverse the child nodes and accumulate the variable names. That is the case, for example, for Form. Again, the only method implementation that does something different is Var, which returns the x argument.

```
@Algebra
public interface QLAlg<E, S, F> {
  F Form(String name, List<S> body);
  S If(E cond, S then);
  S Question(String name,String label,String type);
  E Lit(int n);
  E Var(String x);
  E GEq(E lhs, E rhs);
}
```

Figure 3. Object Algebra interface of the MiniQL abstract syntax.

Unlike the standard OOP implementation, Object Algebras support adding operations without changing existing code. For instance, the renaming operation mentioned above could be realized as follows:

```
class Rename<E, S, F> implements QLAlg<E, S, F> {
  QLAlg<E, S, F> alg;
  String from, to;

  F Form(String n, List<S> b) {
    return alg.Form(n, b);
  }
  S Question(String n, String l, String t) {
    n = n.equals(from) ? to : n;
    return alg.Question(n, l, t);
  }
  ...
  E Var(String x) {
    x = x.equals(from) ? to : x;
    return alg.Var(x);
  }
}
```

Each constructor reconstructs a new node using an auxiliary MiniQL algebra `alg`. Almost all the method implementations reconstruct the structure with no changes using the methods of `alg`. For instance, the `Form` method just recreates the form in the algebra `alg`. The other boilerplate cases are omitted for brevity; the full code can be found in Appendix A.1.2. The two exceptions are the methods `Question` and `Var`, where the identifiers with the given name `from` are renamed to `to`.

Although the Object Algebra encoding of MiniQL solves the problem of extensibility, the traversal code still contains boilerplate code. In both `UsedVars` and `Rename`, the only interesting code is in a small number of cases. Ideally, we would like to write only the code for the interesting cases, and somehow "inherit" the tedious traversal code.

3.3 Shy: An Object Algebra Framework for Traversals

To deal with the boilerplate problem we created **Shy**: a Java Object Algebras framework, which provides a number of generic traversals at the cost of a single annotation. The key idea in **Shy** is to automatically create highly generic Object Algebras, which encapsulate common types of traversals. In particular **Shy** supports generic *queries* and *transformations*.

```
class UsedVars implements QLAlgQuery<Set<String>> {
  public Monoid<Set<String>> m() {
    return new SetMonoid<String>();
  }
  public Set<String> Var(String name) {
    return Collections.singleton(name);
  }
}
```

Figure 4. MiniQL used variables, implemented with **Shy**.

```
class Rename<E, S, F> extends QLAlgTrans<E, S, F> {
  String from, to;
  Rename(QLAlg<E,S,F> alg, String from, String to) {
    super(alg);
    this.from = from;
    this.to = to;
  }
  public S Question(String n, String l, String t) {
    n = n.equals(from) ? to : n;
    return qLAlg().Question(n, l, t);
  }
  public E Var(String x) {
    x = x.equals(from) ? to : x;
    return qLAlg().Var(x);
  }
}
```

Figure 5. MiniQL renaming, implemented with **Shy**.

The two types of traversals are, for instance, sufficient to capture the used variables and renaming operations.

With **Shy**, programmers just need to add the `@Algebra` annotation to the definition of `QLAlg` to get the code for generic queries and transformations. An example of that annotation is already shown in Figure 3. Triggered by the annotation, **Shy** generates base traversal interfaces with Java 8 **default** methods which can then be overridden to implement specific behavior. For instance, for the MiniQL algebra, **Shy** generates interfaces `QLAlgQuery` and `QLAlgTrans` which can be used to implement `UsedVars` and `Rename` in only a fraction of the code[2].

The **Shy**-based implementation of both operations is shown in Figure 4 and 5. In contrast to Figure 2, the code in Figure 4 is much shorter. By implementing the `QLAlgQuery` and `QLAlgTrans` interface, only the methods `Question` and `Var` need to be overridden: all the other methods perform basic accumulation for queries and basic reconstruction in the case of transformations. For queries the only extra thing a programmer has to do is to provide an instance of a monoid, which is used to specify how to accumulate the results during the traversal. Similarly, for transformations, the programmer needs to pass an algebra for providing the constructors for creating the result of a transformation.

---

[2] The generated code is available in Appendix A.1.3 and A.1.4.

Client code To use the queries and operations on a questionnaire like the one in Figure 1, we need a function to create a structure using the generic MiniQL interface:

```
<E, S, F> F makeQL(QLAlg<E, S, F> alg) {
  return alg.Form("DriverLicense", Arrays.asList(
    alg.Question("name", "Name?", "string"),
    alg.Question("age", "Age?", "integer"),
    alg.If(alg.GEq(alg.Var("age"),
                   alg.Lit(18)),
            alg.Question("license",
              "License?", "boolean")))));
}
```

Since both queries and transformations are implementations of the MiniQL interface, they can be passed to the `makeQL` function defined above:

```
println(makeQL(new UsedVars()));
println(makeQL(new Rename<>(new UsedVars(),
    "age", "AGE")));
```

This code prints out `[age]` and `[AGE]`, which are the set of used variables before and after renaming, respectively. Note how the `Rename` transformation transforms the questionnaire into the `UsedVars` algebra.

The remainder of the paper provides the details and implementation techniques used in **Shy**. Besides basic queries and transformations, **Shy** also supports two generalizations of these types of traversals called *generalized queries* and *contextual transformations*.

4. Queries

This section shows the ideas behind generic queries and how they are implemented in **Shy**. A query is an operation that traverses a structure and computes some aggregate value. The inspiration for queries comes from similar types of traversals used in functional programming libraries, such as "Scrap your Boilerplate" [11].

The following code shows a variant of the type of expressions shown in Section 2, represented as the Object Algebra interface `ExpAlg`.

```
@Algebra
interface ExpAlg<Exp> {
  Exp Var(String s);
  Exp Lit(int i);
  Exp Add(Exp e1, Exp e2);
}
```

We will use this minimal Object Algebra interface throughout the rest of the paper to illustrate the various different types of traversals supported by **Shy**. Three different kinds of nodes exist: a numeric literal, a variable or the addition of two expressions. Queries are illustrated by implementing an operation to compute the free variables in an expression.

```
interface FreeVars extends ExpAlg<Set<String>> {
  default Set<String> Var(String s) {
    return Collections.singleton(s);
  }
  default Set<String> Lit(int i) {
    return Collections.emptySet();
  }
  default Set<String> Add(Set<String> e1,
    Set<String> e2) {
    return Stream.concat(e1.stream(), e2.stream())
      .collect(Collectors.toSet());
  }
}
```

Figure 6. Free variables as an Object Algebra.

4.1 Boilerplate Queries

Figure 6 shows a standard approach for computing free variables using Object Algebras[3]. A set of strings is used to collect the names of the free variables. The `Var` method returns a singleton set of `s`, whereas the `Lit` method returns an empty set. The more interesting case is in the `Add` method, where the two sets are joined into one.

The typical pattern of a query is to collect some information from some of the nodes of the structure, and to aggregate the information that comes from multiple child nodes. For example, in the case of free variables, the strings from the `Var` nodes are collected, and in the `Add` nodes the information from multiple children is merged into a single set.

An important observation about queries is that the code to aggregate information tends to be the same: if we had a subtraction node, the code would be essentially identical to `Add`. Moreover, there are only very few types of nodes that contain relevant information for the query. For nodes that contain information that is not relevant to the query, we simply return a neutral value (such as the empty set in `Lit`). Nonetheless, a programmer has to write this boring boilerplate code handling the traversals. While for the small structure presented here this may not look too daunting, in a large structure with dozens or even hundreds of constructors such code becomes a significant burden.

4.2 Generic Queries

A better approach would be to abstract the generic traversal and accumulation code for queries. This way, when programmers need to implement query operations, they can simply reuse the generic traversal code and focus only on dealing with the nodes that do something interesting.

The code that captures the aggregation and collection of information can be captured by a well-known algebraic structure called a *monoid*. Monoids are commonly used in functional programming for such purposes, but they are per-

---

[3] Here and in the following we will use interfaces with **default**-methods (as introduced in Java 8) to combine queries and transformations using multiple inheritance.

```
interface ExpAlgQuery<Exp> extends ExpAlg<Exp> {
 Monoid<Exp> m();
 default Exp Var(String s) { return m().empty(); }
 default Exp Lit(int i) { return m().empty(); }
 default Exp Add(Exp e1, Exp e2) {
   return m().join(e1, e2);
 }
}
```

Figure 7. Generic queries using a monoid.

haps less commonly known in object-oriented programming. The interface of a monoid is defined as follows:

```
   interface Monoid<R> {
     R join(R x, R y);
     R empty();
   }
```

Intuitively, the `join()` method is used to combine the information from substructures, and `empty()` is an indicator of "no information". Using the monoid operations alone, it is possible to write a generic query. Figure 7 shows how this is achieved. In nodes that contain child nodes, such as `Add`, the information is aggregated using `join`. In nodes that contain other information, such as `Var` and `Lit`, the query returns `empty`. This allows concrete queries to be implemented by overriding methods from multiple, different algebras.

4.3 Free Variables with Generic Queries

The `ExpAlgQuery` interface provides an alternative way to define the free variables operation. Instead of directly defining the free variables operation, `ExpAlgQuery` can be inherited, provided that the method `m()` is implemented. In the case of free variables, the monoid returned by `m()` is an implementation of the `Monoid` interface for sets:

```
class SetMonoid<X> implements Monoid<Set<X>> {
 public Set<X> empty() {
   return Collections.emptySet();
 }
 public Set<X> join(Set<X> x, Set<X> y) {
   Set<X> tmp = new HashSet<>(x);
   tmp.addAll(y);
   return tmp;
 }
}
```

The method `empty()` corresponds to the empty set, and `join()` is implemented as union. Using this monoid the free variables operation is defined as follows:

```
interface FreeVars extends ExpAlgQuery<Set<String>>
{
 default Monoid<Set<String>> m() {
   return new SetMonoid<String>();
 }
 default Set<String> Var(String s) {
   return Collections.singleton(s);
 }
}
```

There are two important differences to the implementation in Figure 6. The first difference is that the monoid to be used needs to be specified. However, the code for the monoid is still quite short (only requires two method implementations) and, more importantly, it is highly reusable. Indeed the `SetMonoid` is reused throughout the paper on various examples of queries. Because monoid instances are so general purpose, the **Shy** library already contains many common monoid implementations. Users do not usually have to define these instances themselves. The second difference is that now only the case for variables needs to be defined: the other cases are inherited from `ExpAlgQuery`.

The traversal code in `ExpAlgQuery` is entirely mechanical and can be automatically generated. This is precisely what **Shy** does. Annotating algebra interfaces, such as `ExpAlg`, with the annotation `@Algebra`, triggers automatic generation of generic query interfaces, such as `ExpAlgQuery`. The general template in **Shy** for an algebra $Alg<X_1, ..., X_n>$, with constructors $f_1, ..., f_m$ is shown next:

```
interface AlgQ<R> extends Alg<R,...,R> {
 Monoid<R> m();

 default R fᵢ() {
   return m().empty();
 }

 default R fⱼ(R p₁, ..., R pₖ) {
   return m().join(p₁, m().join(p₂, ..., m().join(
     pₖ₋₁, pₖ)...)));
 }
 ...
}
```

Note that interface Alg_Q extends `Alg` so that all type parameters are unified as type `R`. All arguments to a constructor f_j are combined with `join` from the monoid `m()`. Arguments with primitive types, like **int**, **boolean** or `String`, are ignored by default.

5. Generalized Queries

The previous section introduced simple queries where each constructor contributes to a single monoid. Recursive data types, however, often have multiple syntactic categories, for instance expressions and statements. In such multi-sorted Object Algebras each sort is represented by a different type parameter in the algebra interface. In this section we present *generalized queries*, where each such type parameter can be instantiated to different monoids. It turns out that, for some operations, this generalized version of queries is needed.

Example: Data Dependencies. A simple example of a generalized query is the extraction of the data dependencies between assignment statements and variables in simple imperative programs. To express this query, the simple `ExpAlg` is first extended with statements using the `StatAlg` interface defined as follows:

```
interface G_StatAlgQuery<Exp, Stat>
  extends StatAlg<Exp, Stat> {

 Monoid<Exp> mExp();
 Monoid<Stat> mStat();

 default Stat Assign(String x, Exp e) {
  return mStat().empty();
 }
 default Stat Seq(Stat s1, Stat s2) {
  return mStat().join(s1, s2);
 }
}
```

Figure 8. Default implementation of generalized queries over many-sorted statement algebra.

```
@Algebra
public interface StatAlg<Exp, Stat> {
 Stat Seq(Stat s1, Stat s2);
 Stat Assign(String x, Exp e);
}
```

The `StatAlg` interface defines statement constructors for sequential composition (`Seq`) and assignment (`Assign`). The interface does not extend the `ExpAlg` interface; we rely on multiple inheritance of Java interfaces to combine implementations of these interfaces later (see Figure 9).

The generated default implementation of queries over statements is shown in Figure 8, while the generated code for expressions (`G_ExpAlgQuery`) is presented in Appendix A.1.5. Note that the interface declares two monoids, one for each sort. Since the `Assign` and `Seq` constructors create statements, they return elements of the `mStat()` monoid. Furthermore, because it is impossible to automatically join a monoid over one type with a monoid over another type, the e argument in `Assign` is ignored. As a result, a concrete implementation normally has to override this case to deal with the transition from expressions to statements.

Data dependencies are created by assignment statements: for a statement `Assign(String x, Exp e)` method, the variable x will depend on all variables appearing in e. The result of extracting such dependencies can be represented as binary relation (a set of pairs). In expressions we need to collect the free variables, which can be stored in a set of strings. Thus in this traversal two monoids are involved: a monoid for a set of pairs of strings; and a monoid for a set of strings.

To implement the extraction of data dependencies only two cases have to be implemented: the variable (`Var`) case from the `ExpAlg` signature; and the assignment (`Assign`) case from the `StatAlg` signature. The implementation is shown in Figure 9. Note that the `Assign` case takes the input `Set<String>` e and uses it to create the dependency relation. The propagation of dependencies across sequential composition is automatic, as is the propagation of the set of variables through the different types of expressions.

```
interface DepGraph extends
  G_ExpAlgQuery<Set<String>>,
  G_StatAlgQuery<Set<String>, Set<Pair<String,
       String>>> {
 default Monoid<Set<String>> mExp(){
  return new SetMonoid<>();
 }
 default Monoid<Set<Pair<String, String>>> mStat(){
  return new SetMonoid<>();
 }
 default Set<String> Var(String x){
  return Collections.singleton(x);
 }
 default Set<Pair<String, String>>
       Assign(String x, Set<String> e){
  Set<Pair<String, String>> deps = new HashSet<>();
  e.forEach(y -> deps.add(new Pair<>(x, y)));
  return deps;
 }
}
```

Figure 9. Dependency graph with a generalized query.

Client code A structure using the generic interfaces `ExpAlg` and `StatAlg` is created as follows:

```
<E, S, A extends ExpAlg<E> & StatAlg<E, S>>
 S makeStat(A a) {
  return a.Seq(
   a.Assign("x", a.Add(a.Var("x"), a.Lit(3))),
   a.Assign("y", a.Add(a.Var("x"), a.Var("z"))));
}
```

Note that here the argument of `makeStat` must implement both `ExpAlg` and `StatAlg`. To achieve this in Java, `makeStat` has a type parameter A which is required to implement both interfaces. Using `makeStat` we can pass an instance of `DepGraph` to compute the dependencies of the statement.

```
println(makeStat(new DepGraph(){}));
```

The result is [<x, x>, <y, x>, <y, z>], as expected.

6. Transformations

Queries are a way to extract information from a data structure. Transformations, on the other hand, allow data structures to be transformed into new structures. Just as with queries, we can distinguish code that deals with traversing the data structure from code that actually transforms the structure. In this section we show how to avoid most traversal boilerplate code in the context of transformations using **Shy**.

6.1 Transformations, Object Algebra Style

A simple example of a transformation algebra, using the Object Algebra interface `ExpAlg`, is substituting expressions for variables. A manual implementation based on Object Algebras is shown in Figure 10.

The expression to be substituted, and the variable to substitute for are provided by the methods `e()` and `x()` respectively. The method `expAlg()` is an instance of `ExpAlg` on

```
interface SubstVar<Exp> extends ExpAlg<Exp> {
  ExpAlg<Exp> expAlg();
  String x(); Exp e();
  default Exp Var(String s) {
    return s.equals(x())? e(): expAlg().Var(s);
  }
  default Exp Lit(int i) { return expAlg().Lit(i); }
  default Exp Add(Exp e1, Exp e2) {
    return expAlg().Add(e1, e2);
  }
}
```

Figure 10. A normal algebra-based implementation of variable substitution.

which the transformation is based. Since Object Algebras are factories, the transformation is executed immediately during construction of tree structures. For instance, calling `Var("x")` on a `SubstVar` object with `x()` returning `"x"` immediately returns the result of `e()` (the original variable expression is never created). In the other cases, the original structure is recreated in the algebra `expAlg()`.

The following code shows how the transformation could be used:

```
FreeVars fv = new FreeVars() {};
SubstVar<Set<String>> subst =
  new SubstVar<Set<String>>() {
    public ExpAlg<Set<String>> expAlg() {
      return fv;
    }
    public String x() { return "x"; }
    public Set<String> e() {
      return fv.Add(fv.Lit(1),fv.Var("y"));
    }
};
Set<String> res = subst.Var("x");
```

The `SubstVar` interface is instantiated with `expAlg()` returning an instance of the `FreeVars` algebra defined earlier (e.g., Figure 6). The `x()` method returns the variable to be substituted (`"x"`). Finally, the `e()` returns a new expression $1 + y$ over the `fv` algebra. When expressions are created on the `subst` algebra, the result is the set of free variables *after* the substitution has taken place. As a result, `res` will contain only `"y"`.

Note that this allows pipelining of transformations: there is no reason `expAlg()` cannot return yet another transformation algebra, for instance, a another instance of `SubstVar` realizing a different substitution. We elaborate on composing transformations this way in Section 11.2.

Unfortunately, we again observe the problem of traversal-only boilerplate code: the `Lit` and `Add` methods of Figure 10 simply delegate to the base algebra `expAlg()`, without doing any real work.

6.2 Generic Traversal Code

The boilerplate code in transformations can be avoided by creating a super-interface containing default methods per-

```
interface ExpAlgTransform<Exp> extends ExpAlg<Exp> {
  ExpAlg<Exp> expAlg();
  default Exp Var(String s) {
    return expAlg().Var(s);
  }
  default Exp Lit(int i) { return expAlg().Lit(i); }
  default Exp Add(Exp e1, Exp e2) {
    return expAlg().Add(e1, e2);
  }
}
```

Figure 11. Traversal-only base interface for implementing transformations of expressions.

```
interface Alg_T<X_1,…,X_n> extends Alg<X_1,…,X_n> {
  Alg<X_1, ..., X_n> alg();

  default X_i f_j(X_p^1 p_1, …, X_p^k p_k) {
    return alg().f_j(p_1, …, p_k);
  }
  …
}
```

Figure 12. Generic template for generating boilerplate of transformations.

forming the traversal (shown in Figure 11). A concrete transformation can then selectively override the cases of interest. Variable substitution can now be implemented as follows:

```
interface SubstVar<Exp>
    extends ExpAlgTransform<Exp> {
  String x(); Exp e();
  default Exp Var(String s) {
    return s.equals(x())? e(): expAlg().Var(s);
  }
}
```

In this case, only the method `Var()` is overridden.

Just like in the case of queries, the traversal code in `ExpAlgTransform` is entirely mechanical and can be automatically generated by **Shy**. Figure 12 shows the general template for the generated code. Here `Alg_T` extends `Alg` with the same type parameters and the base algebra `alg()` is declared inside.

7. Contextual Transformations

The previous section introduced a simple template for defining transformations. Transformations in this style may only depend on global context information (e.g., `x()`, `e()`). Many transformations, however, require context information that might change during the traversal itself. In this section we instantiate algebras over function types to obtain transformations which pass information down during traversal. Instead of having the algebra methods delegate directly to base algebra (e.g., `expAlg()`), this now happens indirectly through closures that propagate the context information.

Figure 13 shows the general template for an `Alg<X1, X2, ..., Xn>`, with constructors `f1, ..., fm`. Note that inter-

```
interface Alg_CT<C,X_1,…,X_n>
   extends Alg<Function<C,X_1>,…,Function<C,X_n>> {
 Alg<X_1, …, X_n> alg();

 default Function<C, X_i> f_j(Function<C,X_p^1> p_1, …,
     Function<C,X_p^k> p_k) {
  return (c) -> alg().f_j(p_1.apply(c), …, p_k.apply(c
     ));
 }
 …
}
```

Figure 13. Generic template for generating boilerplate of contextual transformations.

face Alg_{CT} extends `Alg` and instantiates the type parameters to `Function`s from the context argument `C` to the corresponding sort X_i. Each constructor method now creates an anonymous function which, when invoked, calls the functions received as parameters (p_1 to p_k) and only then creates a structure over the `alg()` algebra.

7.1 Example: Conversion to De Bruijn Indices

An example of a contextual transformation is converting variables to De Bruijn indices in the lambda calculus [6]. Using De Bruijn indices, a variable occurrence is identified by a natural number equal to the number of lambda terms between the variable occurrence and its binding lambda term. Lambda terms expressed using De Bruijn indices are useful because they are invariant with respect to alpha conversion.

The conversion to De Bruijn indices uses an object algebra interface `LamAlg` with constructors for lambda abstraction (`Lam`) and application (`Apply`). See below:

```
@Algebra
public interface LamAlg<Exp> {
 Exp Lam(String x, Exp e);
 Exp Apply(Exp e1, Exp e2);
}
```

Shy automatically generates traversal code for transformation for both `LamAlg` and `ExpAlg`: `G_LamAlgTransform` and `G_ExpAlgTransform`, respectively. The generated transformation code can be found in Appendix A.1.6. Using these interfaces, the conversion to De Bruijn indices is realized as shown in Figure 14. Note again that only the relevant cases are overridden: `Var` (from `ExpAlg`) and `Lam` (from `LamAlg`).

Client code   A structure using the generic interfaces `ExpAlg` and `LamAlg` is created as follows:

```
<E, A extends ExpAlg<E> & LamAlg<E>>
   E makeLamExp(A alg) {
 return alg.Lam("x", alg.Lam("y",
    alg.Add(alg.Var("x"), alg.Var("y"))));
}
```

It simply generates "$\lambda x.\lambda y.\ x + y$" as a generic lambda expression. By instantiating the `DeBruijn` interface with a `PrintExpLam` algebra (shown in Appendix A.1.8) passed in as the base algebra, we can write the client code as follows:

```
interface DeBruijn<E> extends
  G_ExpAlgTransform<List<String>, E>,
  G_LamAlgTransform<List<String>, E> {
 default Function<List<String>, E> Var(String p0) {
  return xs ->
     expAlg().Var("" + (xs.indexOf(p0) + 1));
 }

 default Function<List<String>, E> Lam(String x,
     Function<List<String>, E> e) {
  return xs ->
     lamAlg().Lam("", e.apply(cons(x, xs)));
 }
}
```

Figure 14. Converting variables to De Bruijn indices.

```
DeBruijn<String> deBruijn = new DeBruijn<String>() {
 PrintExpLam alg = new PrintExpLam();
 public ExpAlg<String> expAlg() { return alg; }
 public LamAlg<String> lamAlg() { return alg; }
};
println(makeLamExp(deBruijn)
  .apply(Collections.emptyList()));
```

The printed output is "`\.\.(2 + 1)`", which corresponds to original lambda term, but without variables.

8. Desugaring Transformations

In Section 6, we presented transformations, as well as the generic traversals generated by **Shy**. Although different constructors can be used in transforming a data structure, the generic transformations generated by **Shy** are type-preserving: they transform structures over one type (e.g., expressions) to different structures in the same type.

Desugaring transformations eliminate syntactic constructs by transforming them to a combination of constructs in a smaller base language. In this section we describe how **Shy** can be applied to implement compositional desugarings in a type-safe and extensible manner. In particular, the type system will enforce that the resulting language is indeed "smaller", and that consequently the desugared construct is guaranteed to be fully eliminated.

As an example, consider extending `ExpAlg` with a doubling construct which multiplies its argument expression by two. The Object Algebra interface that implements doubling is defined as follows:

```
@Algebra
public interface DoubleAlg<E> {
 E Double(E e);
}
```

An expression `Double(e)` can be desugared to `Add(e,e)`. The following code realizes this transformation by extending the `ExpAlgTransform` interface, generated by **Shy**:

```
interface Desugar<E> extends DoubleAlg<E>,
    ExpAlgTransform<E> {
 default E Double(E e) {
  return expAlg().Add(e, e);
```

```
  }
}
```

The interface `Desugar` exports the language `DoubleAlg` and `ExpAlg`, but `expAlg()` (which is used as a factory for output expressions) has type `ExpAlg`. Since `ExpAlg` does not contain `Double`, the `Double` constructor cannot be used to construct the output. As a result, the algebra `Desugar` transforms into is guaranteed to be without any occurrences of `Double`.

Expressions are created over the combined languages `DoubleAlg` and `ExpAlg`:

```
<E, Alg extends DoubleAlg<E> & ExpAlg<E>>
 E makeExp(Alg a) {
   return a.Add(a.Lit(5), a.Double(a.Var("a")));
 }
```

To illustrate the use of the `Desugar` algebra, here is the code to print an expression in desugared form:

```
   ExpAlg<String> print = new PrintExp();
   Desugar<String> desugar = new Desugar<String>() {
     @Override
     public ExpAlg<String> expAlg() {
       return print;
     }
   };
   System.out.println(makeExp(desugar));
```

The kind of desugarings presented in this section are limited to bottom-up, compositional desugaring, corresponding to factory methods directly invoking methods in a different algebra. As a result, these desugarings are executed in a bottom-up fashion: the arguments are always desugared before an expression itself is transformed. Because the transformations are generic with respect to the carrier object of the argument (as indicated by the type parameter `E`) it is impossible to look at the arguments. This prevents desugarings to perform complex, deep pattern matching on the argument structure. An added benefit, however, is that the desugaring is automatically deforested: intermediate expression trees are never created.

9. Extensible Queries and Transformations

Shy queries and transformations inherit modular extensibility from the Object Algebra design pattern. New transformations or queries are simply added by extending the interfaces generated by **Shy**. More interestingly, however, it is also possible to extend the data type with new constructors. Here we briefly describe how queries and transformations can be extended in this case.

9.1 Linear Extensibility

Consider again the extension of the expression language with lambda and application constructs (cf. Section 7). This requires changing the free variables query, since variables bound by `Lam` expressions need to be subtracted from the set of free variables of the body. Instead of reimplementing the query from scratch, it is possible to modularly extend the existing `FreeVars` query:

```
interface FreeVarsWithLambda extends FreeVars,
    LamAlgQuery<Set<String>> {
  default Set<String> Lam(String x,Set<String> f) {
    return f.stream().filter(y -> !y.equals(x))
           .collect(toSet());
  }
}
```

The interface `FreeVarsWithLambda` extends both the original `FreeVars` query and the base query implementation that was generated for the `LamAlg` interface defining the language extension. Note again, that only the relevant method (`Lam`) needs to be overridden.

For transformations the pattern is similar. To illustrate extension of transformation, consider the simple transformation that makes all variable occurrences unique, to distinguish multiple occurrences of the same name:

```
interface Unique<E> extends ExpAlgTransform<E> {
  int nextInt();
  default E Var(String s) {
    return expAlg().Var(s + nextInt());
  }
}
```

The `Unique` transformation uses a helper method `nextInt` which returns consecutive integers on each call. The basic transformation simply renames `Var` expressions. If, again, the expression language is extended with lambda constructs, the transformation needs to be updated as well to make the variable in the binding position of lambda expression unique. The following code shows how this can be done in a modular fashion:

```
interface UniqueWithLambda<E> extends Unique<E>,
    LamAlgTransform<E> {
  default E Lam(String x, E e) {
    return lamAlg().Lam(x + nextInt(), e);
  }
}
```

Note that the transformation uses the `lamAlg()` algebra (from `LamAlgTransform`), to create lambda expressions.

Figure 15 and 16 give a high level overview of query and transformation extension using the examples for `FreeVars` and `Unique`, respectively. In the case of queries, the abstract `m()` method will be shared by both the `FreeVars` and `FreeVarsWithLambda` interfaces. On the other hand, transformations are based on multiple base algebras, for sets of data type constructors (e.g., `expAlg()` and `lamAlg()`).

Note finally that, in the current implementation of **Shy** transformations, it is assumed that the language signatures `ExpAlg` and `LamAlg` are completely independent. This is however, not an essential requirement. An alternative design could have `LamAlg` be a proper extension of `ExpAlg` (i.e. `LamAlg<E> extends ExpAlg<E>`). In that case, the generated `LamAlgTransform` would need to refine the return type of the `expAlg()` method.
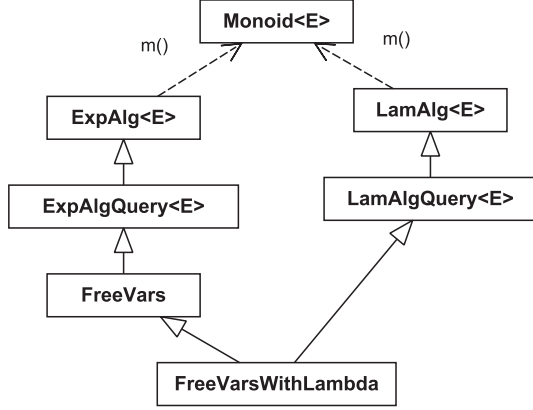
**Figure 15.** Extension of the `FreeVars` query.

Figure 16. Extension of the `Unique` transformation.

9.2 Independent Extensibility

Both Figure 15 and Figure 16 show extensions of queries and transformations where the extensions directly inherit from the concrete implementation of the operations (i.e. `FreeVars` and `Unique`, respectively). It is also possible to make the extensions more independent using multiple inheritance of interfaces with default methods.

For instance, the `Unique` transformation for the `LamAlg` language could also have been implemented independently of `Unique`:

```
interface UniqueLam<E> extends LamAlgTransform<E> {
  int nextInt();
  default E Lam(String x, E e) {
    return lamAlg().Lam(x + nextInt(), e);
  }
}
```

Note that this interface declares the required dependency on `nextInt()`. Both transformations are combined independently through multiple inheritance of interfaces. Both `Unique` and `UniqueLam` declare the method `nextInt()`, but since the declarations are abstract, they are identified. Imple-

mentations of the combined interface need to provide concrete implementations for `nextInt()`, `expAlg()` and `lamAlg()`:

```
class Combine<E> implements Unique<E>,UniqueLam<E> {
  public int nextInt() { ... }
  public ExpAlg<E> expAlg() { ... }
  public LamAlg<E> lamAlg() { ... }
}
```

For queries the pattern is the same, except that only a concrete implementation of `m()` has to be provided for the combined interface.

10. Shy Implementation

Shy is implemented using the standard Java annotations framework (`javax.annotation`) packaged in the Java Development Kit. All of the generic traversals are automatically generated by **Shy** for Object Algebra interfaces annotated with `@Algebra`. When an Object Algebra interface is annotated with `@Algebra`, **Shy** retrieves the required information from the interface (such as the names and types of factory methods) using reflection. **Shy** then generates the code based on the templates shown earlier. Furthermore, the **Shy** framework includes the `Monoid` interface as well as several useful implementations of it.

A major advantage of using standard Java annotations is that the code generation of the generic traversals can be done transparently: users do not need to use or install a tool to generate that code. As a result **Shy** is as simple to use as a conventional library. With little configuration effort, the code generation is automatically enabled in IDEs like Eclipse or IntelliJ. Finally, the framework is very small (around 885 source lines of code), so it can be easily be customized, if needed.

11. Case Study

To illustrate the utility of **Shy** we have implemented a number of queries and transformations in the context of QL, a DSL for questionnaires which has been implemented using Object Algebras before [9]. QL is similar to MiniQL, except that it additionally features an if-then-else construct, computed questions (which will appear read only), and a richer expression language. For more information on the features of QL we refer to [7].

11.1 QL Queries and Transformations

The queries extract derived information from a QL program, such as the set of used variables, the data and control dependencies between questions, and the global type environment. The transformations include two transformations of language extensions to the base language. The first realizes a simple desugaring of "unless(c)..." to "if(not(c))...". The second desugaring statically unfolds a constant bound loop construct ("repeat (i)...") and renames variables occurring below it accordingly. Finally, we have implemented a sim-

Operation	Exp (18)	Stmt (5)	Form (1)	%
Collect variables	1			4%
Data dependencies		3	1	17%
Control dependencies		4	1	21%
Type environment		2		8%
Rename variable	1	2		13%
Inline conditions		4		17%
Desugar "unless"		1		4%
Desugar "repeat"	1	3		16%

Table 1. Number of overriden cases per query and transformation in the context of the QL implementation.

ple rename variable operation, and a flattening normalizer which inlines the conditions of nested if-then constructs.

Table 1 shows the number of cases that had to be overridden to implement each particular operation. The top row shows the number of constructs for each syntactic category in QL (Exp, Stmt, and Form). As can be seen, none of the operations required implementing all cases. The last column shows the number of overridden cases as a percentage. For this set of queries and transformations, almost no expression cases needed to be overridden, except the "Var" case in collect variables, rename variable and desugar "repeat"[4]. The cases required for desugaring include the case of the language extension (e.g. `Unless` and `Repeat`, respectively). These cases are not counted in the total in the first row but are used to compute the percentage.

11.2 Chaining Transformations

A typical compiler consists of many transformations chained together in a pipeline. **Shy** transformations support this pattern by passing transformation algebras as the base algebra to the implementation of another transformation. For instance, the desugar unless transformation desugars the "unless" statement to "if" statements in another algebra. The latter can represent yet another transformation.

In the context of QL, "unless" desugaring, condition inlining and variable renaming can be chained together as follows:

```
alg = new Desugar<>(new Inline<>(
  new Rename<>(Collections.singletonMap("x", "y"),
    new Format()))));
```

The chained transformation `alg` first desugars "unless", then inlines conditions, and finally renames xs to ys. The `Rename` transformation gets as base algebra an instance of `Format`, a pretty printer for QL.

The algebra `alg` can now be used to create questionnaires:

```
Function<IFormatWithPrecedence, IFormat> pp
  = alg.form("myForm", Arrays.asList(
      alg.unless(alg.var("x"),
      alg.question("x", "X?", new TBoolean())))));
```

---

[4] Note, however, that the dependency extraction queries reuse the collect variables query on expressions.


Figure 17. Performance comparison of control dependencies query.

Since inlining is a contextual transformation, the result of constructing this simple questionnaire is a function object representing the "to be inlined" representation of the questionnaire after desugaring. The `IFormatWithPrecedence` and `IFormat` types are formatting operations, respectively representing expressions and statements; these types originate from the `Format` algebra passed to `Rename`.

Calling the function with a boolean expression representing **true** will trigger inlining of conditions and renaming. The result is then a formatting object (`IFormat`) which can be used to print out the transformed questionnaire:

```
form myForm { if (true && !y) y "X?" boolean }
```

As can be seen, the variable x has been renamed to y. The (renamed) condition y is now negated, because of the desugaring of "unless". Finally, the result of inlining conditions can be observed from the conjunction in the **if** statement.

11.3 Shy Performance vs Vanilla ASTs

We compared the performance characteristics of the operations implemented using **Shy** with respect to vanilla implementations based on ordinary AST classes with ordinary methods representing the transformations and queries. In the vanilla implementation, the program was parsed into an AST structure, and then the operation was invoked and measured. In the case of the **Shy** queries, constructing the "AST" corresponds to executing the query, so we measured that. For context-dependent transformations, however, building the "AST" corresponds to constructing the function to execute the transformation, hence we only measured invoking this function. The vanilla query implementations use the same monoid structures as in **Shy**.

The operations were executed on progressively larger QL programs (up to 140Kb). The QL programs represent questionnaires describing a binary search problem (a number

Figure 18. Performance comparison of inline conditions transformation.

Component	SLOC
Object Algebra interfaces	71
Collect variables	10
Data dependencies	30
Control dependencies	58
Type environment	15
Rename variable	27
Inline conditions	48
Desugar "unless"	10
Desugar "repeat"	35
Total Shy-based operations	**304**
AST based implementation	**661**

Table 2. Source Lines of Code (SLOC) statistics: **Shy** implementation vs Vanilla AST implementation.

guessing game) and are automatically generated, with increasing search spaces. The benchmarks were executed on a 2.6GHz MacBook Pro Intel Core i5 with 8GB memory. The JVM was run in 64bit server mode and was given 4Gb of heapspace to minimize the effect of garbage collection pauses. Each benchmark was run without measuring first, to warm-up the JVM. The measurements presented here do not include warm-up time.

The comparison of the control dependencies query is shown in Figure 17. The plot shows that the performance is quite comparable. On average, the **Shy** implementation of the query seems a little slower. This is probably caused by the extensive use of interfaces in the **Shy** framework, whereas the AST-based implementation only uses abstract and concrete classes. For transformations the performance difference is slightly more pronounced. Figure 18 shows the performance comparison of the inline conditions transformation. The greater difference can be explained by the fact that creating a new structure in a **Shy** transformation involves dynamically dispatched method calls instead of statically bound constructor calls.

11.4 Shy vs Vanilla Regarding Code Size

The percentages shown in Table 1 illustrate the structure-shyness of queries and transformations implemented using **Shy**. Table 2 shows the absolute number of source lines of code (SLOC, lines of code without counting empty lines and comments). For the vanilla AST-based implementation we only show the total SLOC count, since all operations are scattered over the respective AST classes. For the **Shy** implementation, each query and operation is realized a separate class or interface, extending the interfaces generated by **Shy**. In total, one can observe that the **Shy**-based implementation requires less than half of the number of lines of code

required in the vanilla implementation. Note also that **Shy** supports a much more modular design, where both the AST data type and the set of operations can be extended without having to modify existing code.

12. Related Work

Structure-shy traversals have been an active research topic. Our approach to structure shy traversal is unique in that it supports separate compilation, modular type checking and data type extension. Furthermore, it can be applied in mainstream languages such as Java. While some approaches support some of these features, to the best of our knowledge, no approach supports all of them.

Adaptive Object-Oriented Programming (AOOP) AOOP is an extension of object-oriented programming aimed at increasing the flexibility and maintainability of programs [17]. AOOP promotes the idea of structure-shyness to achieve those goals. In AOOP it is possible to select parts of a structure that should be visited. This is useful to do traversals on complex structures and focus only on the interesting parts of the structure relevant for computing the final output. The original approach to AOOP was based on a domain-specific language [17]. DJ is an implementation of AOOP in Java using reflection [22]. More recently DemeterF [5] improved previous approaches to AOOP by providing support for type-safe traversals, generics and data-generic function generation. **Shy** shares with AOOP the use of structure-shyness as a means to increase flexibility and adaptability of programs. Most AOOP approaches, however, are not type-safe. The exception is DemeterF where a custom type system was designed to ensure type-safety of generic functions. Unlike DemeterF, which is a separate language, **Shy** is a Java library. Moreover, compilation of DemeterF programs is implemented through static weaving, and thus appears to preclude separate compilation.

Strategic Programming Strategic programming is an approach to data structure traversal, which originated in term rewriting. Visser et al. extended the rewriting strategies of [3] with generic one-level traversal operators [27] enabling a style of term rewriting where computations are represented by simple, conditional rewrite rules, but the application of such rules is controlled separately using the concept of a strategy. Strategies can be primitive (e.g., "fail") or composed using combinators (e.g., "try *s* else *s'*"). These and other combinators were formalized in a core language for strategic rewriting in [26].

The strategy concept has since then been ported to other paradigms. JJTRAVELER is an OO framework for strategic programming [28]. Lämmel et al. introduced typed strategy combinators in Haskell [14]. The relation between strategic programming and AOOP has been explored in [16].

The key tenet of strategic programming is separation of concerns: actual computation and traversal are specified separately. In **Shy**, the traversal of a data structure is also specified separately (in a super-interface), however, it is fixed for specific styles of queries and transformations. For instance, both queries and transformations employ an innermost, bottom-up strategy.

The distinction between queries and transformations also originates from existing work in strategic programming. Lämmel et al. [14] discuss type unifying and type preserving traversals. Type unifying traversals correspond to queries, where all data type constructors are unified into a single monoid. Analogously, **Shy** transformations are type preserving in the sense that a transformation is an algebra which maps constructor calls to another algebra of the same type.

The ASF+SDF program transformation system distinguishes transforming and accumulating traversals, which correspond to our transformations and queries, respectively [25]. Furthermore, an *accumulating transforming* traversal combines both styles, by tupling the accumulated result and the transformed tree. This combination could easily be generated by **Shy** by having the boilerplate code construct the monoid and the transformed term in parallel (see Fig. 19).

Structure-Shy Traversals with Visitors A standard way to remove boilerplate in OOP is to use *default visitors* [18]. Default visitors can be used in similar ways to our generic traversals. Many programmers using the visitor pattern create such default visitor implementations to avoid boilerplate code. There are two important differences to our work. Firstly conventional visitors are not extensible in a type-safe way. Secondly, with **Shy** the code for generic traversals is automatically generated, whereas with default visitors such code usually has to be implemented by hand.

Visser [28] adapted the strategy combinators of Stratego [26, 27] to combinators that operate on object-oriented Visitors [8]. The resulting framework JJTRAVELER solves the problem of entangling traversal control within the `accept`

```
interface AccuTrafoExp<M,E> extends
    ExpAlg<Pair<M,E>> {
  Monoid<M> m(); ExpAlg<E> expAlg();
  default Pair<M,E> Var(String s) {
    return new Pair<>(m().empty(),
      expAlg().Var(s));
  }
  default Pair<M,E> Lit(int i) {
    return new Pair<>(m().empty(),
      expAlg().Lit(i));
  }
  default Pair<M,E> Add(Pair<M,E> e1, Pair<M,E> e2)
  {
    return new Pair<>(m().join(e1.fst(), e2.fst()),
      expAlg().Add(e1.snd(), e2.snd())));
  }
}
```

Figure 19. Combining query and transformation.

methods, or in the Visitors themselves (which only allow static specialization). A challenge not addressed by JJTRAVELER is type safety of traversal code: either the combinators needs to be redefined for each data type, or client code needs to cast the generic objects of type `AnyVisitable` to the specific type. Even if specific combinators would be generated, however, the traversed data types would not be extensible.

Another approach to improve upon the standard Visitor pattern is presented in Palsberg et al. [23]. This work particularly addressed the fact that traditional Visitors operate on a fixed set of classes. As a result, the data type can not be extended without changing all existing Visitors as well. The proposed solution is a generic `Walkabout` class which accesses sub-components of arbitrary data structures using reflection. Unfortunately, the heavy use of reflection make `Walkabout`s significantly slower than traditional Visitors. The authors state that the `Walkabout` class could be generated to improve performance, but note that the addition of a class could trigger regeneration. As a result the pattern does not support separate compilation. Our solution obtains the same kind of default behavior for traversal, without losing extensibility, type safety, or separate compilation.

Whereas the `Walkabout` provides generic navigation over an object structure, this navigation can be programmed explicitly using *guides* [4]. Guides insert one level of indirection between recursing on the children of a node in `visit` methods: the guide decides how to proceed the traversal. Since guides needs to define how to proceed for each type that will be visited, they suffer from the same extensibility problem as ordinary Visitors. Generic guides, on the other hand, are dynamically typed and use reflection to call appropriate `visit` methods. The `Walkabout` can be formulated as such a generic guide.

Structure-Shy Traversals in Functional Programming In functional programming there has been a lot of research on type-safe structure-shy traversals. Lämmel and Peyton

Jones' "Scrap your Boilerplate" (SyB) [11–13] series introduced a practical design pattern for doing generic traversals in Haskell. The simple queries and transformations in **Shy** were partly inspired by SyB. However SyB and **Shy** use very different implementation techniques. SyB is implemented in Haskell and relies on a run-time type-safe cast mechanism. This approach allows SyB traversals to be encoded once-and-forall using a single higher-order function called `gfoldl`. In contrast, in **Shy** Java annotations are used to generate generic traversals for each structure.

A drawback of SyB traversals is that they are notoriously slow, partly due to the use of the run-time cast [1]. Another notable difference between SyB and **Shy** is with respect to extensibility. While **Shy** supports extensibility of both traversals and structures, the original SyB approach did not support any extensibility. Only in later work, Lämmel and Peyton Jones proposed an alternative design for SyB, based on type classes [30]. This design supports extensibility of traversals, but not of the traversed structures.

Closest to **Shy** is an approach proposed by Lämmel et al. [15] for dealing with the so-called "*large bananas*". A large banana corresponds to the fold algebra of a complex structure. Object Algebras, which we use in our work, are an OO encoding of fold algebras [19, 20]. However Lämmel et al. work has not dealt with extensibility. Interestingly in their future work Lämmel et al. did mention that they would like "to cope with incomplete or extensible systems of datatypes".

Language Extensions for Queries Inspired by XPath/X-Query there has been some work on adding support for similar types of queries on object-oriented structures. For example the work on $C\omega$ [2] extends C# with generalized member access, which allows simple XPath-like path expressions. Thus in $C\omega$ it is possible to express queries quite concisely. However, in contrast to **Shy** $C\omega$ is a language extension and it does not deal with transformations.

Eliminating Boilerplate in Design Patterns Design patterns [8] improve the design and modularity of object-oriented programs. However, the implementation of design patterns sometimes requires significant amounts of boilerplate code. There has been some work on implementing design patterns in AspectJ to achieve reusability and modularity [10], and thus eliminate some of the boilerplate code. A challenge with traditional design patterns, however, is that the boilerplate code is not always mechanical, due to many possible implentation choices. Our work proposes a number of design patterns for traversals. Because implementing these design patterns by hand would be quite tedious, we automatically generate the code for such design patterns using **Shy**. Fortunatelly, in contrast to many of the traditional design patterns, the code for generic traversals is highly regular and easy to generate, and can be completely eliminated.

Object Algebras **Shy** traversals are based on Object Algebras [19]. The original motivation for object algebras was as a design pattern for OO programming that allowed improved extensibility and modularity of programs. Using object algebras it is possible to solve the well-known "Expression Problem" [29]. Later work [21, 24] has explored the use of *Object Algebra combinators*, and generalizations of object algebras to improve expressiveness and modularity. In particular it has been claimed that Object Algebras can be used to do *feature-oriented programming* [21], and to encode *attribute grammars* [24]. One domain where Object Algebras are especially useful is in the implementation of (extensible) languages. The QL language used in our case study is based on Gouseti et al. [9]. That work provides a realistic implementation of an extensible *domain-specific language* using Object Algebras. In contrast to our work, which focus on removing boilerplate code, previous work on Object Algebras was mostly motivated by improved programming support for extensibility and modularity. The combination of extensibility (inherited for free from Object Algebras) and structure-shy type-safe traversals adds a new dimension to our work that, as far as we know, has not been explored previously.

13. Conclusion

This paper showed how various types of default traversals for complex structures can be automatically provided by **Shy**. **Shy** traversals are written directly in Java and are type-safe, extensible and separately compilable. There has always been a tension between the correctness guarantees of static typing, and the flexibility of untyped/dynamically-typed approaches. **Shy** shows that even in type systems like Java's, it is possible to get considerable flexibility and adaptability for the problem of boilerplate code in traversals of complex structures, without giving up modular static typing.

There are many avenues for future work. One area of research is to extend **Shy** traversals to support flexible traversal strategies, similarly to strategic programming [3, 25, 26]. Another line of work worth exploring is to adopt generalizations of object algebras [21] for added expressiveness of **Shy** traversals.

Acknowledgments

References

[1] M. D. Adams and T. M. DuBuisson. Template your boilerplate: Using Template Haskell for efficient generic programming. In *Proceedings of the 2012 ACM SIGPLAN Haskell symposium*, Haskell '12, pages 13–24, 2012.

[2] G. M. Bierman, E. Meijer, and W. Schulte. The essence of data access in comega. In *ECOOP*, 2005.

[3] P. Borovanskỳ, C. Kirchner, H. Kirchner, P.-E. Moreau, and M. Vittek. ELAN: A logical framework based on computational systems. *Electronic Notes in Theoretical Computer Science*, 4:35–50, 1996.

[4] M. Bravenboer and E. Visser. Guiding visitors: Separating navigation from computation. Technical Report UU-CS-2001-42, Institute of Information and Computing Sciences, Utrecht University, 2001.

[5] B. Chadwick and K. Lieberherr. Weaving generic programming and traversal performance. In *AOSD'10*, 2010.

[6] N. de Bruijn. Lambda calculus notation with nameless dummies, a tool for automatic formula manipulation, with application to the Church-Rosser theorem. *Indagationes Mathematicae (Proceedings)*, 75(5):381–392, 1972.

[7] S. Erdweg, T. van der Storm, M. Völter, M. Boersma, R. Bosman, W. R. Cook, A. Gerritsen, A. Hulshout, S. Kelly, A. Loh, et al. The state of the art in language workbenches. In *Software Language Engineering*. Springer, 2013.

[8] E. Gamma, R. Helm, R. Johnson, and J. Vlissides. *Design Patterns: Elements of Reusable Object-oriented Software*. Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 1995. ISBN 0-201-63361-2.

[9] M. Gouseti, C. Peters, and T. v. d. Storm. Extensible language implementation with Object Algebras (short paper). In *GPCE'14*, 2014.

[10] J. Hannemann and G. Kiczales. Design pattern implementation in Java and aspectJ. In *OOPSLA '02*, 2002.

[11] R. Lammel and S. P. Jones. Scrap your boilerplate: A practical design pattern for generic programming. In *TLDI'03*, 2003.

[12] R. Lämmel and S. P. Jones. Scrap more boilerplate: Reflection, zips, and generalised casts. In *ICFP '04*, 2004.

[13] R. Lämmel and S. P. Jones. Scrap your boilerplate with class: Extensible generic functions. In *ICFP '05*, 2005.

[14] R. Lämmel and J. Visser. Typed combinators for generic traversal. In *Practical Aspects of Declarative Languages*, pages 137–154. Springer, 2002.

[15] R. Lämmel, J. Visser, and J. Kort. Dealing with large bananas. In J. Jeuring, editor, *Workshop on Generic Programming*, Ponte de Lima, July 2000. Technical Report UU-CS-2000-19, Universiteit Utrecht.

[16] R. Lämmel, E. Visser, and J. Visser. Strategic programming meets adaptive programming. In *AOSD '03*, 2003.

[17] K. J. Lieberherr. *Adaptive Object Oriented Software: The Demeter Method with Propagation Patterns*. PWS Publishing, 1996.

[18] M. E. Nordberg. Variations on the visitor pattern, 1996.

[19] B. C. d. S. Oliveira and W. R. Cook. Extensibility for the masses, practical extensibility with object algebras. In *ECOOP '12*, 2012.

[20] B. C. d. S. Oliveira, M. Wang, and J. Gibbons. The visitor pattern as a reusable, generic, type-safe component. In *OOPSLA'08*, 2008.

[21] B. C. d. S. Oliveira, T. van der Storm, A. Loh, and W. R. Cook. Feature-oriented programming with object algebras. In *ECOOP '13*, 2013.

[22] D. Orleans and K. J. Lieberherr. DJ: Dynamic adaptive programming in Java. In *Reflection 2001*. Springer-Verlag, 2001.

[23] J. Palsberg and C. B. Jay. The essence of the visitor pattern. In *COMPSAC'98*, 1998.

[24] T. Rendel, J. I. Brachthäuser, and K. Ostermann. From object algebras to attribute grammars. In *OOPSLA '14*, 2014.

[25] M. G. J. van den Brand, P. Klint, and J. J. Vinju. Term rewriting with traversal functions. *ACM Trans. Softw. Eng. Methodol.*, 12(2):152–190, Apr. 2003.

[26] E. Visser and Z.-e.-A. Benaissa. A core language for rewriting. *Electronic Notes in Theoretical Computer Science*, 15: 422–441, 1998.

[27] E. Visser, Z.-e.-A. Benaissa, and A. Tolmach. Building program optimizers with rewriting strategies. In *ICFP'98*, 1998.

[28] J. Visser. Visitor combination and traversal control. In *OOPSLA '01*, 2001.

[29] P. Wadler. The Expression Problem. Email, Nov. 1998. Discussion on the Java Genericity mailing list.

[30] P. Wadler and S. Blott. How to make ad-hoc polymorphism less ad hoc. In *POPL '89*, 1989.

A. Appendix

A.1 Complete Code

A.1.1 OO Approach for `usedVars` and `rename`

Below is the complete code for Fig. 2 (left). It implements `usedVars` and `rename` in the QL example, as an OO approach.

```java
class Form {
  String name;
  List<Stmt> body;
  Form(String id, List<Stmt> body) {
    this.name = id;
    this.body = new ArrayList<Stmt>(body);
  }
  Set<String> usedVars() {
    Set<String> vars = new HashSet<>();
    body.forEach(s -> vars.addAll(s.usedVars()));
    return vars;
  }
  Form rename(String n1, String n2) {
    List<Stmt> ss = new ArrayList<>();
    for (Stmt s: body) ss.add(s.rename(n1, n2));
    return new Form(name, ss);
  }
}

abstract class Stmt {
  abstract Set<String> usedVars();
  abstract Stmt rename(String n1, String n2);
}

class If extends Stmt {
  Exp cond;
  Stmt then;
  If(Exp cond, Stmt then) {
```

```
    this.cond = cond;
    this.then = then;
  }
  Set<String> usedVars() {
    Set<String> vars =
       new HashSet<>(cond.usedVars());
    vars.addAll(then.usedVars());
    return vars;
  }
  If rename(String n1, String n2) {
    return new If(cond.rename(n1, n2),
       then.rename(n1, n2));
  }
}


class Question extends Stmt {
  String name, label, type;
  Question(String n, String l, String t) {
    this.name = n;
    this.label = l;
    this.type = t;
  }
  Set<String> usedVars() {
    return emptySet();
  }
  Question rename(String n1, String n2) {
    String newN = name.equals(n1) ? n2 : name;
    return new Question(newN, label, type);
  }
}


abstract class Exp {
  abstract Set<String> usedVars();
  abstract Exp rename(String n1, String n2);
}


class Lit extends Exp {
  int n;
  Lit(int n) {
    this.n = n;
  }
  Set<String> usedVars() {
    return emptySet();
  }
  Lit rename(String n1, String n2) {
    return new Lit(n);
  }
}


class Var extends Exp {
  String x;
  Var(String name) {
    this.x = name;
  }
  Set<String> usedVars() {
    return Collections.singleton(x);
  }
  Var rename(String n1, String n2) {
    String newN = x.equals(n1) ? n2 : x;
    return new Var(newN);
  }
}


class GEq extends Exp {
  Exp lhs, rhs;
```

```
  GEq(Exp lhs, Exp rhs) {
    this.lhs = lhs;
    this.rhs = rhs;
  }
  Set<String> usedVars() {
    Set<String> vars = new HashSet<>(lhs.usedVars());
    vars.addAll(rhs.usedVars());
    return vars;
  }
  GEq rename(String n1, String n2) {
    return new GEq(lhs.rename(n1, n2),
       rhs.rename(n1, n2));
  }
}
```

A.1.2 Rename implementing the QLAlg interface

The following code gives the implementation of Rename that implements QLAlg in Section 3.2.

```
class Rename<E, S, F> implements QLAlg<E, S, F> {
  private QLAlg<E, S, F> alg;
  private String from, to;
  public Rename(QLAlg<E, S, F> alg, String from,
      String to) {
    this.alg = alg;
    this.from = from;
    this.to = to;
  }
  public F Form(String id, List<S> stmts) {
    return alg.Form(id, stmts);
  }
  public S If(E c, S t) {
    return alg.If(c, t);
  }
  public S Question(String n, String l, String t) {
    n = n.equals(from) ? to : n;
    return alg.Question(n, l, t);
  }
  public E Lit(int n) {
    return alg.Lit(n);
  }
  public E Var(String x) {
    x = x.equals(from) ? to : x;
    return alg.Var(x);
  }
  public E GEq(E l, E r) {
    return alg.GEq(l, r);
  }
}
```

A.1.3 QLAlgQuery: generated code

The generated code for QLAlgQuery by **Shy** in Fig. 4.

```
public interface QLAlgQuery<R>
    extends QLAlg<R, R, R> {

  Monoid<R> m();

  default R Form(java.lang.String p0,
     java.util.List<R> p1) {
    R res = m().empty();
    res = m().join(res, m().fold(p1));
    return res;
  }
```

```java
  default R Geq(R p0, R p1) {
   R res = m().empty();
   res = m().join(res, p0);
   res = m().join(res, p1);
   return res;
  }

  default R If(R p0, R p1) {
   R res = m().empty();
   res = m().join(res, p0);
   res = m().join(res, p1);
   return res;
  }

  default R Lit(int p0) {
   R res = m().empty();
   return res;
  }

  default R Question(java.lang.String p0,
    java.lang.String p1, java.lang.String p2) {
   R res = m().empty();
   return res;
  }

  default R Var(java.lang.String p0) {
   R res = m().empty();
   return res;
  }

}
```

A.1.4 `QLAlgTransform` and `QLAlgTrans`: generated code

The code for `QLAlgTransform` and its class representation `QLAlgTrans` for use, generated by **Shy**. See Fig. 5.

```java
public interface QLAlgTransform<A0, A1, A2>
    extends QLAlg<A0, A1, A2> {

 QLAlg<A0, A1, A2> qLAlg();

 default A2 Form(java.lang.String p0,
    java.util.List<A1> p1) {
   return qLAlg().Form(p0, p1);
 }

 default A0 Geq(A0 p0, A0 p1) {
   return qLAlg().Geq(p0, p1);
 }

 default A1 If(A0 p0, A1 p1) {
   return qLAlg().If(p0, p1);
 }

 default A0 Lit(int p0) {
   return qLAlg().Lit(p0);
 }

 default A1 Question(java.lang.String p0,
    java.lang.String p1, java.lang.String p2) {
   return qLAlg().Question(p0, p1, p2);
 }

 default A0 Var(java.lang.String p0) {
```

```java
   return qLAlg().Var(p0);
 }

}

public class QLAlgTrans<A0, A1, A2>
    implements QLAlgTransform<A0, A1, A2> {

 private QLAlg<A0, A1, A2> alg;

 public QLAlgTrans(QLAlg<A0, A1, A2> alg) {
   this.alg = alg;
 }

 public QLAlg<A0, A1, A2> qLAlg() {return alg;}

}
```

A.1.5 `G_ExpAlgQuery`: generated code

The generated code for `G_ExpAlgQuery` by **Shy** in Fig. 9.

```java
public interface G_ExpAlgQuery<A0>
    extends ExpAlg<A0> {

 Monoid<A0> mExp();

 default A0 Add(A0 p0, A0 p1) {
   A0 res = mExp().empty();
   res = mExp().join(res, p0);
   res = mExp().join(res, p1);
   return res;
 }

 default A0 Lit(int p0) {
   A0 res = mExp().empty();
   return res;
 }

 default A0 Var(java.lang.String p0) {
   A0 res = mExp().empty();
   return res;
 }

}
```

A.1.6 `G_ExpAlgTransform` and `G_LamAlgTransform`: generated code

Below is the generated code for `G_ExpAlgTransform` and `G_LamAlgTransform` by **Shy** in Fig. 14.

```java
public interface G_ExpAlgTransform<A, B0>
    extends ExpAlg<Function<A, B0>> {

 ExpAlg<B0> expAlg();

 default <B> List<B> substListExpAlg(List<Function<
    A, B>> list, A acc) {
   List<B> res = new ArrayList<B>();
   for (Function<A, B> i : list)
    res.add(i.apply(acc));
   return res;
 }
```

```java
  default Function<A, B0> Add(Function<A, B0> p0,
      Function<A, B0> p1) {
    return acc -> expAlg().Add(p0.apply(acc),
      p1.apply(acc));
  }

  default Function<A, B0> Lit(int p0) {
    return acc -> expAlg().Lit(p0);
  }

  default Function<A, B0> Var(java.lang.String p0) {
    return acc -> expAlg().Var(p0);
  }

}

public interface G_LamAlgTransform<A, B0>
    extends LamAlg<Function<A, B0>> {

  LamAlg<B0> lamAlg();

  default <B> List<B> substListLamAlg(List<Function<
      A, B>> list, A acc) {
    List<B> res = new ArrayList<B>();
    for (Function<A, B> i : list)
      res.add(i.apply(acc));
    return res;
  }

  default Function<A, B0> Apply(Function<A, B0> p0,
      Function<A, B0> p1) {
    return acc -> lamAlg().Apply(p0.apply(acc),
      p1.apply(acc));
  }

  default Function<A, B0> Lam(java.lang.String p0,
      Function<A, B0> p1) {
```

```java
    return acc -> lamAlg().Lam(p0, p1.apply(acc));
  }

}
```

A.1.7 `Util.cons`: an auxiliary method

The auxiliary method `Util.cons` is implemented as follows, for the De Bruijn example in Fig. 14.

```java
public class Util {
  public static <X> List<X> cons(X x, List<X> l) {
    l = new ArrayList<>(l);
    l.add(0, x);
    return l;
  }
}
```

A.1.8 `PrintExpLam`: a pretty printer for `ExpAlg` and `LamAlg`

The class `PrintExpLam` used in Section 7.1.

```java
class PrintExpLam implements ExpAlg<String>,
    LamAlg<String> {
  public String Lam(String x, String e) {
    return "\\" + x + "." + e;
  }
  public String Apply(String e1, String e2) {
    return "(" + e1 + " " + e2 + ")";
  }
  public String Var(String s) { return s; }
  public String Lit(int i) { return i + ""; }
  public String Add(String e1, String e2) {
    return "(" + e1 + " + " + e2 + ")";
  }
}
```