

Scrap your Boilerplate with Object Algebras

Name1

Affiliation1

Email1

Name2 Name3

Affiliation2/3

Email2/3

Abstract

Traversing complex *Abstract Syntax Trees* (ASTs) typically requires large amounts of tedious boilerplate code. For many operations most of the code simply walks the structure, and only a small portion of the code implements the functionality that motivated the traversal in the first place. This paper presents a type-safe Java framework called **Shy** that removes much of this boilerplate code. In **Shy** *Object Algebras* are used to describe complex and extensible AST structures. Using Java annotations **Shy** generates generic boilerplate code for various types of traversals. For a concrete traversal, users of **Shy** can then inherit from the generated code and override only the interesting cases. Consequently, the amount of code that users need to write is significantly smaller. Moreover, traversals using the **Shy** framework are also much more *structure shy*, becoming more adaptive to future changes or extensions to the AST structure. To prove the effectiveness of the approach, we applied **Shy** in the implementation of a domain-specific questionnaire language. Our results show that for a large number of traversals there was a significant reduction in the amount of user-defined code.

1. Introduction

Various language processing tools or libraries for programming languages, domain-specific languages, mark-up languages like HTML, or data-interchange languages like XML or JSON require complex *Abstract Syntax Tree* (AST) structures. In those applications ASTs are the key data structure needed to model the various constructs of these languages. Such ASTs have various different types of nodes, which can range from a few dozen to several hundred kinds of nodes (for example in the ASTs of languages like Java or Cobol).

Static types are helpful to deal with such complex ASTs. Static types formalize the distinction between different kinds of nodes. Furthermore the distinctions are helpful to ensure that traversals over these ASTs have an appropriate piece of code that deals with each different type of node. This can prevent a large class of run-time errors that would not otherwise be detected.

Unfortunately, when traversing such ASTs, the number of nodes and the enforced type distinctions between nodes can lead to so-called *boilerplate code* [11]: code that is similar

for most types of nodes and which essentially just walks the structure. Operations where the “interesting” code is limited to only a small portion of nodes are called *structure shy* [17]. A typical example is computing the free variables of an expression in some programming language. In this case, the interesting code occurs in the nodes representing the binding and variable constructs. In all other cases, the code would just deal with walking the structure. In ASTs with dozens or hundreds of kinds of nodes, having to explicitly write code for each kind of node is both tedious and error-prone.

The boilerplate problem in implementing traversals has received considerable attention in the past. For example, both *Adaptive Object-Oriented Programming* (AOOP) [17] and *Strategic Programming* [3, 25] are aimed partly at solving this problem. Most approaches to AOOP and strategic programming use some meta-programming techniques, such as code generation or reflection. The use of meta-programming offers programmers an easy way to avoid having to write boilerplate code. This has important benefits: users have to write much less code; and the code becomes much more adaptive to changes. **However Nevertheless**, such meta-programming based approaches usually come at the cost of other desirable properties, such as type-safety, extensibility or separate compilation. The functional programming community has also studied the problem before. For instance, the popular “*Scrap your boilerplate*” [11] approach supports type-safety and separate compilation. **However But** most of the techniques used in functional languages cannot be easily ported to mainstream OO languages like Java, and are limited in terms of extensibility.

This paper presents a Java framework called **Shy** that allows users to define type-safe and extensible structure-shy operations. **Shy** uses *Object Algebras* [18] to describe ASTs, and to write operations over ASTs in a style similar to writing folds in functional programming. However unlike standard folds in functional programming, Object Algebras are *extensible*: when new kinds of nodes are introduced in the data type, the operations can be extended as well, without changing existing code.

In **Shy** Object Algebra interfaces are combined with Java annotations to generate generic traversal code automatically. The generated code accounts for four different types of traversals: *queries*; *transformations*; *generalized queries*;

and *contextual transformations*. Each of these four types of traversals is implemented as an Object Algebra. Programmers who want to implement structure-shy traversals can then inherit from one of these four Object Algebras, and override only the cases that deal with the interesting parts of the traversal. Consequently traversals written in **Shy** are:

- **Small in size:** With **Shy** the amount of code that programmers need to write a structure-shy traversal is significantly smaller. Often traversals in **Shy** are implementable in just a few lines of code, even for complex ASTs with hundreds of different types of nodes.
- **Adaptive and structure shy:** Traversals written with **Shy** can omit boilerplate code, making traversals more adaptive to future changes or extensions to the data type.
- **Type-safe:** **Shy** traversals are directly written in Java and the Java type-system ensures type-safety. No run-time casts are needed for generic traversal code or for user-defined traversal code.
- **Extensible with separate compilation:** Traversals inherit type-safe extensibility from Object Algebras. Both traversals and the AST structures are extensible. Thus it is possible to reuse traversal code for ASTs extended with additional node types. Furthermore **Shy** traversals support separate compilation.
- **Implemented in plain Java:** **Shy** traversals do not require a new tool or language. The approach is library based and only uses Java annotations.

To prove the effectiveness of the approach, we have applied **Shy** in the implementation of QL, a domain-specific language (DSL) for defining questionnaires that has been used before as a benchmark language [7, 9]. Our results (see details in Section 10) show that for a large number of traversals there was a significant reduction in the amount of user-defined code: only 4% to 21% of the AST cases had to be implemented in comparison with code written without **Shy**.

Although **Shy**'s functional programming inspired idioms are new to mainstream Java programmers, **its basis on standard Java (?) mechanisms** makes **Shy** easy to implement, integrate in existing environments, and use. Moreover, although Java was chosen as the implementation language for **Shy**, our approach should apply to any object-oriented language with support for generics and annotations.

In summary, the contributions of this paper are:

- **Design patterns for generic traversals.** We provide a set of design patterns for various types of traversals using Object Algebras, including: *queries*, *transformations*, *generalized queries* and *contextual transformations*.
- **The Shy Java framework.** We show that those design patterns can be automatically generated for a given Object Algebra interface. The **Shy** framework¹ realizes this

idea by using Java annotations to automatically generate generic traversals.

- **Case study.** We illustrate the benefits of **Shy** using a case study based on the QL domain-specific language. The results of our case study show significant savings in terms of user-defined traversal code. The case study also shows that **Shy** does not incur significant performance overhead compared to a regular AST-based implementation.

2. Background: Object Algebras

Object Algebras capture a design pattern to solve the Expression Problem [28]. As a result Object Algebras support modular and type-safe extensibility in two dimensions: data variants and operations. Object Algebras can be used conveniently to model recursive data types, such as ASTs. Here we briefly introduce Object Algebras using the example of a simple expression language.

```
interface ExpAlg<E> {
    E Lit(int n);
    E Add(E l, E r);
}
```

The above code shows an *Object Algebra interface* which models literal and addition expressions. The interface resembles an abstract factory interface, but instead of constructing objects of a concrete type, the result is abstract through the type parameter E. A concrete Object Algebra will implement such an interface by instantiating the type parameter with a type that constructs objects for a particular purpose. For instance, the class Eval (which implements evaluation of **expressions**) is as follows:

```
class Eval implements ExpAlg<Integer> {
    public Integer Lit(int n) {
        return n;
    }
    public Integer Add(Integer l, Integer r) {
        return l + r;
    }
}
```

Using this algebra to create expressions will immediately evaluate the expression to its result. To add another operation the interface ExpAlg is implemented again. As an example, the class Print presented below shows a simple printing operation on expressions. Both examples illustrate extension in the dimension of operations.

```
class Print implements ExpAlg<String> {
    public String Lit(int n) { return "" + n; }
    public String Add(String l, String r) {
        return l + " + " + r;
    }
}
```

The other dimension of extensibility (adding data variants) is realized by extending the Object Algebra interface itself. For instance, the interface MulAlg defined below extends the original ExpAlg with multiplication expressions.

¹ **Note to Reviewers:** Due to the anonymous review process we submit a file bundle with the implementation together with the paper submission. If the paper is accepted we will make the code publicly available.

```
interface MulAlg<E> extends ExpAlg<E> {
    E Mul(E l, E r);
}
```

The key feature of Object Algebras is that the existing operations (evaluation and printing) can now be extended to support multiplication without having to change the existing Eval and Print classes. For example, evaluation is extended with multiplication as follows:

```
class MulEval extends Eval
implements MulAlg<Integer> {
    public Integer Mul(Integer l, Integer r) {
        return l * r;
    }
}
```

In summary, Object Algebra interfaces describe recursive data types, and implementations of those interfaces represent operations. Extending the interface allows developers to extend the data type with new data variants, and (re)implementing the interface allows developers to define operations. Finally, extension in both dimensions is fully type-safe and does not compromise separate compilation. Although for space reasons this section does not show how to use Object Algebras, such examples will be provided throughout the paper.

3. An Overview of Shy

This section starts by illustrating the problem of boilerplate code when implementing traversals of complex structures. It then shows how **Shy** addresses the problem using a combination of Object Algebras [18] and Java annotations.

3.1 Traversing Object-Oriented ASTs

We start by introducing the problem of boilerplate code by considering a simplified variant of the QL language used in our case study [9], called MiniQL. Just like QL, MiniQL can be used to describe interactive questionnaires. An example is shown in Figure 1. The questionnaire first asks for the user's name and age, and then, if the age is greater than or equal to 18, asks if the user has a driver's license. Because of the conditional construct, the last question only will appear when the user is actually eligible to have driver's license.

MiniQL's abstract syntax contains forms, statements (if-then and question) and expressions (only literals, variables and greater-than-or-equal). A traditional OO implementation is shown on the left Figure 3. A form (class Form) has a name and consists of a list of statements. Statements are conditionals (If) which contain an expression and a statement body, and questions (Question) which have a name, label and type. Expressions are standard, but limited to literals (Lit), variables (Var) and greater-than-or-equal (GEq).

The code in Figure 3 also shows a query over MiniQL structures, namely the collection of used variables. The operation is defined using the method `usedVars`, declared in the abstract superclasses Stmt and Exp (omitted for brevity),

```
form DriverLicense {
    name: "What is your name?" string
    age: "What is your age?" integer
    if (age >= 18)
        license: "Have a driver's license?" boolean
}
```

Figure 1. Example QL questionnaire: driver's license.

```
@Algebra
public interface QLAlg<E, S, F> {
    F Form(String name, List<S> body);
    S If(E cond, S then);
    S Question(String name, String label, String type);
    E Lit(int n);
    E Var(String x);
    E GEq(E lhs, E rhs);
}
```

Figure 2. Object Algebra interface of the MiniQL abstract syntax. (Switch places with Fig.3, TODO.)

and implemented in the concrete statement and expression classes. As can be seen, the only interesting bit of code is the `usedVars` method in class Var. All other implementations merely deal with aggregating results of their child nodes, or returning a default empty set.

The boilerplate code exhibited in the `usedVars` query often also applies to transformations. Consider for example a rename transformation which takes a Form and returns another form where **all the specified** variables are renamed. Again, the only interesting cases would be in the Var and Question classes, where the actual renaming is applied. All other classes, however, require boilerplate to recreate the structure. The full code in Appendix A.1.1 contains a simple example of such a rename operation as well.

In addition to the significant amount of boilerplate code, there is another drawback to the traditional OO solution, which is that it does not support extensibility along the dimension of operations. Each new operation requires pervasive changes across the AST classes.

3.2 Modeling MiniQL with Object Algebras

The right-hand side of Figure 3 shows the used variables operation implemented using Object Algebras. The operation is a class implementing the Object Algebra interface (QLAlg) shown in Figure 2.

The UsedVars class provides an implementation for each of the methods in the object algebra interface, which together define the full used variables operation. Since the result of collecting those variables is `Set<String>`, all the type parameters are set to that type. Most of the method implementations simply traverse the child nodes and accumulate the variable names. That is the case, for example, for Form. Again, the only method implementation that does something different is Var, which returns the `x` argument.

Unlike the standard OOP implementation, Object Algebras support adding operations without changing existing

```

class Form {
  String name; List<Stmt> body;
  Set<String> usedVars() {
    Set<String> vars = new HashSet<>();
    body.forEach(s -> vars.addAll(s.usedVars()));
    return vars;
  }
}

class If extends Stmt {
  Exp cond; Stmt then;
  Set<String> usedVars() {
    Set<String> vars=new HashSet<>(cond.usedVars());
    vars.addAll(then.usedVars());
    return vars;
  }
}

class Question extends Stmt {
  String name, label, type;
  Set<String> usedVars() { return emptySet(); }
}

class Lit extends Exp {
  int n;
  Set<String> usedVars() { return emptySet(); }
}

class Var extends Exp {
  String x;
  Set<String> usedVars() { return singleton(x); }
}

class GEq extends Exp {
  Exp lhs, rhs;
  Set<String> usedVars() {
    Set<String> vars = new HashSet<>{lhs.usedVars()};
    vars.addAll(rhs.usedVars());
    return vars;
  }
}

```

```

class UsedVars implements
  QALg<Set<String>, Set<String>, Set<String>> {

  Set<String> Form(String n, List<Set<String>> b) {
    Set<String> vars = new HashSet<>();
    b.forEach(s -> vars.addAll(s));
    return vars;
  }

  Set<String> If(Set<String> c, Set<String> t) {
    Set<String> vars = new HashSet<>(c);
    vars.addAll(t);
    return vars;
  }

  Set<String> Question(String n,String l,String t) {
    return Collections.emptySet();
  }

  Set<String> Lit(int x) {
    return Collections.emptySet();
  }

  Set<String> Var(String x) {
    return Collections.singleton(x);
  }

  Set<String> GEq(Set<String> l, Set<String> r) {
    Set<String> vars = new HashSet<>(l);
    vars.addAll(r);
    return vars;
  }
}

```

Figure 3. Implementing the “used variables” operation using traditional ASTs (left) and Object Algebras (right).

code. For instance, the renaming operation mentioned above could be realized as follows: (Renaming updated)

```

class Rename<E, S, F> implements QALg<E, S, F> {
  QALg<E, S, F> alg;
  String n1, n2;

  F Form(String n, List<S> b) {
    return alg.Form(n, b);
  }
  S Question(String n, String l, String t) {
    String newN = n.equals(n1) ? n2 : n;
    return alg.Question(newN, l, t);
  }
  ...
  E Var(String x) {
    String newN = x.equals(n1) ? n2 : x;
    return alg.Var(newN);
  }
}

```

Each constructor reconstructs a new node using an **auxiliary** MiniQL algebra alg. Almost all the method implementations reconstruct the structure with no changes using the methods of alg. For instance, the Form method just recreates the form in the algebra alg. The other boilerplate cases are omitted for brevity; the full code can be found in Appendix A.1.2. The two exceptions are the methods Question and Var, where the identifiers **with the given name (n1) are renamed (to n2) are suffixed with “_”**.

Although the Object Algebra encoding of MiniQL solves the problem of extensibility, the traversal code still contains boilerplate code. In both UsedVars and Rename, the only interesting code is in a small number of cases. Ideally, we would like to write only the code for the interesting cases, and somehow “inherit” the tedious traversal code.

```

class UsedVars implements QLAlgQuery<Set<String>> {
    public Monoid<Set<String>> m() {
        return new SetMonoid<String>();
    }
    public Set<String> Var(String name) {
        return Collections.singleton(name);
    }
}

```

Figure 4. MiniQL used variables, implemented with **Shy**.

```

class Rename<E, S, F> extends QLAlgTrans<E, S, F> {
    String n1, n2;
    public Rename(QLAlg<E, S, F> alg, String n1,
        String n2) {
        super(alg); this.n1 = n1; this.n2 = n2;
    }
    public S Question(String n, String l, String t) {
        String newN = n.equals(n1) ? n2 : n;
        return qLAlg().Question(newN, l, t);
    }
    public E Var(String x) {
        String newN = x.equals(n1) ? n2 : x;
        return qLAlg().Var(newN);
    }
}

```

Figure 5. MiniQL renaming, implemented with **Shy**. (Renaming updated.)

3.3 Shy: An Object Algebra Framework for Traversals

To deal with the boilerplate problem we created **Shy**: a Java Object Algebras framework, which provides a number of generic traversals at the cost of a single annotation. The key idea in **Shy** is to automatically create highly generic Object Algebras, which encapsulate common types of traversals. In particular **Shy** supports generic *queries* and *transformations*. The two types of traversals are, for instance, sufficient to capture the used variables and renaming operations.

With **Shy**, programmers just need to add the `@Algebra` annotation to the definition of `QLAlg` to get the code for generic queries and transformations. An example of that annotation is already shown in Figure 2. Triggered by the annotation, **Shy** generates base traversal interfaces with Java 8 **default** methods which can then be overridden to implement specific behavior. For instance, for the MiniQL algebra, **Shy** generates interfaces `QLAlgQuery` and `QLAlgTrans` which can be used to implement `UsedVars` and `Rename` in only a fraction of the code².

The **Shy**-based implementation of both operations is shown in Figure 4 and 5. In contrast to Figure 3, the code in Figure 4 is much shorter. By implementing the `QLAlgQuery` and `QLAlgTrans` interface, only the methods `Question` and `Var` need to be overridden: all the other methods perform basic accumulation for queries and basic identity reconstruc-

tion in the case of transformations. For queries the only extra thing a programmer has to do is to provide an instance of a monoid, which is used to specify how to accumulate the results during the traversal. Similarly, for transformations, the programmer needs to pass an algebra for providing the constructors for creating the result of a transformation.

Client code To use the queries and operations on a questionnaire like the one in Figure 1, we need a function to create a structure using the generic MiniQL interface:

```

<E, S, F> F makeQL(QLAlg<E, S, F> alg) {
    return alg.Form("DriverLicense", Arrays.asList(
        alg.Question("name", "Name?", "string"),
        alg.Question("age", "Age?", "integer"),
        alg.If(alg.GEq(alg.Var("age"),
            alg.Lit(18)),
            alg.Question("license",
                "License?", "boolean"))));
}

```

Since both queries and transformations are implementations of the MiniQL interface, they can be passed to the `makeQL` function defined above:

```

println(makeQL(new UsedVars()));
println(makeQL(new Rename<>(new UsedVars(),
    "age", "AGE")));

```

This code prints out `[age]` and `[AGE]`, which are the set of used variables before and after renaming, respectively. Note how the `Rename` transformation transforms the questionnaire into the `UsedVars` algebra.

The remainder of the paper provides the details and implementation techniques used in **Shy**. Besides basic queries and transformations, **Shy** also supports two generalizations of these types of traversals called *generalized queries* and *contextual transformations*.

4. Queries

This section shows the ideas behind generic queries and how they are implemented in **Shy**. A query is an operation that traverses a structure and computes some aggregate value. The inspiration for queries comes from similar types of traversals used in functional programming libraries, such as “Scrap your Boilerplate” [11].

The following code shows a variant of the type of expressions shown in Section 2, represented as the Object Algebra interface `ExpAlg`.

```

@Algebra
interface ExpAlg<Exp> {
    Exp Var(String s);
    Exp Lit(int i);
    Exp Add(Exp e1, Exp e2);
}

```

We will use this minimal Object Algebra interface throughout the rest of the paper to illustrate the various different types of traversals supported by **Shy**. Three different kinds of nodes exist: a numeric literal, a variable or the addition of

²The generated code is available in Appendix A.1.3 and A.1.4.


```

interface FreeVars extends ExpAlg<Set<String>> {
    default Set<String> Var(String s) {
        return Collections.singleton(s);
    }
    default Set<String> Lit(int i) {
        return Collections.emptySet();
    }
    default Set<String> Add(Set<String> e1,
        Set<String> e2) {
        return Stream.concat(e1.stream(), e2.stream())
            .collect(Collectors.toSet());
    }
}

```

Figure 6. Free variables as an Object Algebra.

two expressions. Queries are illustrated by implementing an operation to compute the free variables in an expression.

4.1 Boilerplate Queries

Figure 6 shows a standard approach for computing free variables using Object Algebras³. A set of strings is used to collect the names of the free variables. The Var method returns a singleton set of *s*, whereas the Lit method returns an empty set. The more interesting case is in the Add method, where the two sets are joined into one.

The typical pattern of a query is to collect some information from some of the nodes of the structure, and to aggregate the information that comes from multiple child nodes. For example, in the case of free variables, the strings from the Var nodes are collected, and in the Add nodes the information from multiple children is merged into a single set.

An important observation about queries is that the code to aggregate information tends to be the same: if we had a subtraction node, the code would be essentially identical to Add. Moreover, there are only very few types of nodes that contain relevant information for the query. For nodes that contain information that is not relevant to the query, we simply return **the** a neutral value (such as the empty set in Lit). Nonetheless, a programmer has to write this boring boilerplate code handling the traversals. While for the small structure presented here this may not look too daunting, in a large structure with dozens or even hundreds of constructors such code becomes a significant burden.

4.2 Generic Queries

A better approach would be to abstract the generic traversal and accumulation code for queries. This way, when programmers need to implement query operations, they can simply reuse the generic traversal code and focus only on dealing with the nodes that do something interesting.

The code that captures the aggregation and collection of information can be captured by a well-known algebraic

³ Here and in the following we will use interfaces with **default**-methods (as introduced in Java 8) to combine queries and transformations using multiple inheritance.

```

interface ExpAlgQuery<Exp> extends ExpAlg<Exp> {
    Monoid<Exp> m();
    default Exp Var(String s) { return m().empty(); }
    default Exp Lit(int i) { return m().empty(); }
    default Exp Add(Exp e1, Exp e2) {
        return m().join(e1, e2);
    }
}

```

Figure 7. Generic queries using a monoid.

structure called a *monoid*. Monoids are commonly used in functional programming for such purposes, but they are perhaps less commonly known in object-oriented programming. The interface of a monoid is defined as follows:

```

interface Monoid<R> {
    R join(R x, R y);
    R empty();
}

```

Intuitively, the join() method is used to combine the information from substructures, and empty() is an indicator of “no information”. Using the monoid operations alone, it is possible to write a generic query. (Double quotes removed) Figure 7 shows how this is achieved. In nodes that contain child nodes, such as Add, the information is aggregated using join. In nodes that contain other information, such as Var and Lit, the query returns empty. This allows concrete queries to be implemented by overriding methods from multiple, different algebras.

4.3 Free Variables with Generic Queries

The ExpAlgQuery interface provides an alternative way to define the free variables operation. Since the result of free variables is a set, the monoid returned by m() is an implementation of the Monoid interface. For this case the following implementation **would suffice**:

```

class SetMonoid<X> implements Monoid<Set<X>> {
    public Set<X> empty() {
        return Collections.emptySet();
    }
    public Set<X> join(Set<X> x, Set<X> y) {
        Set<X> tmp = new HashSet<>(x);
        tmp.addAll(y);
        return tmp;
    }
}

```

The method empty() corresponds to the empty set, and join() is implemented as union. Using this monoid the free variables operation is defined as follows:

```

interface FreeVars extends ExpAlgQuery<Set<String>> {
    default Monoid<Set<String>> m() {
        return new SetMonoid<String>();
    }
    default Set<String> Var(String s) {
        return Collections.singleton(s);
    }
}

```

```
}
}
```

There are two important differences to the implementation in Figure 6. The first difference is that the monoid to be used needs to be specified. However, the code for the monoid is still quite short (only requires two method implementations) and, more importantly, it is highly reusable. Indeed the `SetMonoid` is reused throughout the paper on various examples of queries. Because monoid instances are so general purpose, the **Shy** library already contains many common monoid implementations.^{so} Users do not usually have to define these instances themselves. The second difference is that now only the case for variables needs to be defined: the other cases are inherited from `ExpAlgQuery`.

The traversal code in `ExpAlgQuery` is entirely mechanical and can be automatically generated. This is precisely what **Shy** does. Annotating algebra interfaces, such as `ExpAlg`, with the annotation `@Algebra`, triggers automatic generation of generic query interfaces, such as `ExpAlgQuery`. Figure 8 shows the general template in **Shy** for an algebra `Alg<X1, ..., Xn>`, with constructors `f1, ..., fm`. Note that interface `AlgQ` extends `Alg` so that all type parameters are unified as type `R`. All arguments to a constructor `fj` are combined with `join` from the monoid `m()`. Arguments with primitive types, like `int`, `boolean` or `String`, are ignored by default.

```
interface AlgQ<R> extends Alg<R,...,R> {
    Monoid<R> m();

    default R fj(R p1, ..., R pk) {
        return m().join(p1, m().join(p2, ..., m().join(
            pk-1, pk)...));
    }
    ...
}
```

Figure 8. Generic template for generic queries. (moved)

5. Generalized Queries

The previous section introduced simple queries where each constructor contributes to a single monoid. Recursive data types, however, often have multiple syntactic categories, for instance expressions and statements. In such multi-sorted Object Algebras each sort is represented by a different type parameter in the algebra interface. In this section we present *generalized queries*, where each such type parameter can be instantiated to different monoids. It turns out that, for some operations, this generalized version of queries is needed.

5.1 Example: Data Dependencies

A simple example of a generalized query is the extraction of the data dependencies between assignment statements and variables in simple imperative programs. To express this query, the simple `ExpAlg` is first extended with statements using the `StatAlg` interface defined as follows:

```
interface G_StatAlgQuery<Exp, Stat>
    extends StatAlg<Exp, Stat> {
        Monoid<Exp> mExp(); Monoid<Stat> mStat();
        default Stat Assign(String x, Exp e) {
            return mStat().empty();
        }
        default Stat Seq(Stat s1, Stat s2) {
            return mStat().join(s1, s2);
        }
    }
}
```

Figure 9. Default implementation of generalized queries over many-sorted statement algebra.

```
@Algebra
public interface StatAlg<Exp, Stat> {
    Stat Seq(Stat s1, Stat s2);
    Stat Assign(String x, Exp e);
}
```

At this point one may ask why `StatAlg` is declared independently instead of extending `ExpAlg`. **Shy** supports extensible data structures, but such extensibility is not realized by declaring two Object Algebra interfaces with inheritance. Instead, they are declared separately, and so are the generated generic traversals. A behavior defined by users can then extend both traversals, since Java supports multiple inheritance for interfaces.

The extension `StatAlg` consists of statement constructors for sequential composition (`Seq`) and assignment (`Assign`). The generated default implementation of queries over statements is shown in Figure 9, while the generated code for expressions (`G_ExpAlgQuery`) is presented in Appendix A.1.5. Note that the interface declares two monoids, one for each sort. Since the `Assign` and `Seq` constructors create statements, they return elements of the `mStat()` monoid. Furthermore, because it is impossible to automatically join a monoid over one type with a monoid over another type, the `e` argument in `Assign` is ignored. As a result, a concrete implementation normally has to override this case to deal with the transition from expressions to statements.

Data dependencies are created by assignment statements: for a statement `Assign(String x, Exp e)` method, the variable `x` will depend on all variables appearing in `e`. The result of extracting such dependencies can be represented as binary relation (a set of pairs). In expressions we need to collect the free variables, which can be stored in a set of strings. Thus in this traversal two monoids are involved: a monoid for a set of pairs of strings; and a monoid for a set of strings.

To implement the extraction of data dependencies only two cases have to be implemented: the variable (`Var`) case from the `ExpAlg` signature; and the assignment (`Assign`) case from the `StatAlg` signature. The implementation is shown in Figure 10. Note that the `Assign` case takes the input `Set<String> e` and uses it to create the dependency relation. The propagation of dependencies across sequential composition

is automatic, as is the propagation of the set of variables through the different types of expressions.

```
interface DepGraph extends
    G_ExpAlgQuery<Set<String>>,
    G_StatAlgQuery<Set<String>, Set<Pair<String,
        String>>> {
    default Monoid<Set<String>> mExp(){
        return new SetMonoid<>();
    }
    default Monoid<Set<Pair<String, String>>> mStat(){
        return new SetMonoid<>();
    }
    default Set<String> Var(String x){
        return Collections.singleton(x);
    }
    default Set<Pair<String, String>>
        Assign(String x, Set<String> e){
        Set<Pair<String, String>> deps = new HashSet<>();
        e.forEach(y -> deps.add(new Pair<>(x, y)));
        return deps;
    }
}
```

Figure 10. Dependency graph with a generalized query. (moved)

Client code A structure using the generic interfaces `ExpAlg` and `StatAlg` is created as follows:

```
<E, S, A extends ExpAlg<E> & StatAlg<E, S>>
    S makeStat(A alg) {
    return alg.Seq(
        alg.Assign("x", alg.Add(alg.Var("x"),
            alg.Lit(3))),
        alg.Assign("y", alg.Add(alg.Var("x"),
            alg.Var("z"))));
    }
```

Note that here the argument of `makeStat` must implement both `ExpAlg` and `StatAlg`. To achieve this in Java, `makeStat` has a type parameter `A` which is required to implement both interfaces. Using `makeStat` we can pass an instance of `DepGraph` to compute the dependencies of the statement.

```
println(makeStat(new DepGraph(){}));
```

The result is [`<x, x>`, `<y, x>`, `<y, z>`], as expected.

6. Transformations

Queries are a way to extract information from a data structure. Transformations, on the other hand, allow data structures to be transformed into new structures. Just as with queries, we can distinguish code that deals with traversing the data structure from code that actually transforms the structure. In this section we show how to avoid most traversal boilerplate code in the context of transformations using *Shy*.

6.1 Boilerplate Transformations

A simple example of a transformation algebra, using the Object Algebra interface `ExpAlg`, is substituting expressions

```
interface SubstVar<Exp> extends ExpAlg<Exp> {
    ExpAlg<Exp> expAlg();
    String x(); Exp e();
    default Exp Var(String s) {
        return s.equals(x())? e(): expAlg().Var(s);
    }
    default Exp Lit(int i) { return expAlg().Lit(i); }
    default Exp Add(Exp e1, Exp e2) {
        return expAlg().Add(e1, e2);
    }
}
```

Figure 11. A normal algebra-based implementation of variable substitution.

```
interface ExpAlgTransform<Exp> extends ExpAlg<Exp> {
    ExpAlg<Exp> expAlg();
    default Exp Var(String s) {
        return expAlg().Var(s);
    }
    default Exp Lit(int i) { return expAlg().Lit(i); }
    default Exp Add(Exp e1, Exp e2) {
        return expAlg().Add(e1, e2);
    }
}
```

Figure 12. Traversal-only base interface for implementing transformations of expressions.

for variables. A manual implementation based on Object Algebras is shown in Figure 11.

The expression to be substituted, and the variable to substitute for are provided by the methods `e()` and `x()` respectively. `expAlg()` is an instance of `ExpAlg` on which the transformation is based. Since Object Algebras are factories, the transformation is executed immediately during construction of tree structures. For instance, calling `Var("x")` on a `SubstVar` object with `x()` returning `"x"` immediately returns the result of `e()` (the original variable expression is never created). In the other cases, the original structure is recreated in the algebra `expAlg()`.

Again we observe the problem of traversal-only boilerplate code: the `Lit` and `Add` methods simply delegate to the base algebra `expAlg()`.

6.2 Generic Traversal Code

The boilerplate code in transformations can be avoided by creating a super-interface containing default methods performing the traversal (shown in Figure 12). A concrete transformation can then selectively override the cases of interest. Variable substitution can now be implemented as follows:

```
interface SubstVar<Exp>
    extends ExpAlgTransform<Exp> {
    String x(); Exp e();
    default Exp Var(String s) {
        return s.equals(x())? e(): expAlg().Var(s);
    }
}
```



```

interface AlgT<X1, ..., Xn> extends Alg<X1, ..., Xn> {
    Alg<X1, ..., Xn> alg();

    default Xi fj(Xp1 p1, ..., Xpk pk) {
        return alg().fj(p1, ..., pk);
    }
    ...
}

```

Figure 13. Generic template for generating boilerplate of transformations.

In this case, only the method `Var()` is overridden. To execute substitution, `SubstVar` should be subclassed with implementations for `x()`, `e()`, and `expAlg()`. The base algebra `expAlg()` could for instance be **an algebra for collecting free variables (just `FreeVars` in Section 4), or an algebra for pretty printing the expression with the substitution applied.** Note that this allows pipelining of transformations: there is no reason `expAlg()` cannot return yet another transformation algebra,—we elaborate on composing transformations this way in Section 10.2.

Here we show an example of invoking `SubstVar`. The base algebra is `FreeVars` in Section 4, which collects all the free variables as a set of strings. One can directly create an instance of `SubstVar` as follows:

```

SubstVar<Set<String>> substVar = new SubstVar<Set<
    String>>() {
    public ExpAlg<Set<String>> expAlg() {
        return new FreeVars();
    }
    public String x() {return "x";}
    public Set<String> e() {return expAlg().Var("z");}
};

```

The transformation replaces all variables “x” with “z”. Such an instance `substVar` can now be passed to a structure of `ExpAlg`, which helps to collect the free variables in the expression, before which the transformation is applied. A set of strings is returned as the result.

Just like in the case of queries, the traversal code in `ExpAlgTransform` is entirely mechanical and can be automatically generated by **Shy**. Figure 13 shows the general template for the generated code. Here `AlgT` extends `Alg` with the same type parameters and the base algebra `alg()` is declared inside.

7. Contextual Transformations

The previous section introduced a simple template for defining transformations. Transformations in this style may only depend on global context information (e.g., `x()`, `e()`). Many transformations, however, require context information that might change during the traversal itself. In this section we instantiate algebras over function types to obtain transformations which pass information down during traversal. Instead of having the algebra methods delegate directly to base al-

```

interface AlgCT<C, X1, ..., Xn>
    extends Alg<Function<C, X1>, ..., Function<C, Xn>> {
        Alg<X1, ..., Xn> alg();

        default Function<C, Xi> fj(Function<C, Xp1> p1, ...,
            Function<C, Xpk> pk) {
            return (c) -> alg().fj(p1.apply(c), ..., pk.apply(
                c));
        }
        ...
    }

```

Figure 14. Generic template for generating boilerplate of contextual transformations.

gebra (e.g., `expAlg()`), this now happens indirectly through closures that propagate the context information.

Figure 14 shows the general template for an `Alg<X1, X2, ..., Xn>`, with constructors `f1`, ..., `fm`. Note that interface `AlgCT` extends `Alg` and instantiates the type parameters to Functions from the context argument `C` to the corresponding sort `Xi`. Each constructor method now creates an anonymous function which, when invoked, calls the functions received as parameters (`p1` to `pk`) and only then creates a structure over the `alg()` algebra.

7.1 Example: conversion to De Bruijn indices

An example of a contextual transformation is converting variables to De Bruijn indices in the lambda calculus [6]. Using De Bruijn indices, a variable occurrence is identified by a natural number equal to the number of lambda terms between the variable occurrence and its binding lambda term. Lambda terms expressed using De Bruijn indices are useful because they are invariant with respect to alpha conversion.

The conversion to De Bruijn indices uses an object algebra interface `LamAlg` with constructors for lambda abstraction (`Lam`) and application (`Apply`). See below:

```

@Algebra
public interface LamAlg<Exp> {
    Exp Lam(String x, Exp e);
    Exp Apply(Exp e1, Exp e2);
}

```

Shy automatically generates traversal code for transformation for both `LamAlg` and `ExpAlg`: `G_LamAlgTransform` and `G_ExpAlgTransform`, respectively. For space reasons the generated transformation code is omitted here, but can be found in Appendix A.1.6. Using these interfaces, the conversion to De Bruijn indices is realized as shown in Figure 15. Note again that only the relevant cases are overridden: `Var` (from `ExpAlg`) and `Lam` (from `LamAlg`).

Client code A structure using the generic interfaces `ExpAlg` and `LamAlg` is created as follows:

```

<E, A extends ExpAlg<E> & LamAlg<E>>
    E makeLamExp(A alg) {
        return alg.Lam("x", alg.Lam("y",
            alg.Add(alg.Var("x"), alg.Var("y"))));
    }

```

```

interface DeBruijn<E> extends
    G_ExpAlgTransform<List<String>, E>,
    G_LamAlgTransform<List<String>, E> {
    default Function<List<String>, E> Var(String p0) {
        return xs ->
            expAlg().Var("" + (xs.indexOf(p0) + 1));
    }

    default Function<List<String>, E> Lam(String x,
        Function<List<String>, E> e) {
        return xs ->
            lamAlg().Lam("", e.apply(cons(x, xs)));
    }
}

```

Figure 15. Converting variables to De Bruijn indices.

```

}

```

It simply generates “ $\lambda x.\lambda y. x + y$ ” as a generic lambda expression. By instantiating the DeBruijn interface with a Printer algebra (shown in Appendix A.1.8) passed in as the base algebra, we can write the client code as follows:

```

DeBruijn<String> deBruijn = new DeBruijn<String>() {
    Printer alg = new Printer();
    public ExpAlg<String> expAlg() { return alg; }
    public LamAlg<String> lamAlg() { return alg; }
};
println(makeLamExp(deBruijn)
    .apply(Collections.emptyList()));

```

The printed output is “ $\lambda.\lambda.(2 + 1)$ ”, which corresponds to original lambda term, but without variables.

8. Extensible Queries and Transformations

Shy queries and transformations inherit modular extensibility from the Object Algebra design pattern. New transformations or queries are simply added by extending the interfaces generated by **Shy**. More interestingly, however, it is also possible to extend the data type with new constructors. Here we briefly describe how queries and transformations can be extended in this case.

Consider again the extension of the expression language with lambda and application constructs (cf. Section 7). This requires changing the free variables query, since variables bound by Lam expressions need to be subtracted from the set of free variables of the body. Instead of reimplementing the query from scratch, it is possible to modularly extend the existing FreeVars query:

```

interface FreeVarsWithLambda extends FreeVars,
    LamAlgQuery<Set<String>> {
    default Set<String> Lam(String x, Set<String> f) {
        return f.stream().filter(y -> !y.equals(x))
            .collect(toSet());
    }
}

```

The interface FreeVarsWithLambda extends both the original FreeVars query and the base query implementation that




Figure 16. Extension of the FreeVars query.

was generated for the LamAlg interface defining the language extension. Note again, that only the relevant method (Lam) needs to be overridden.

For transformations the pattern is similar. To illustrate extension of transformation, consider the simple transformation that makes all variable occurrences unique, to distinguish multiple occurrences of the same name:

```

interface Unique<E> extends ExpAlgTransform<E> {
    int nextInt();
    default E Var(String s) {
        return expAlg().Var(s + nextInt());
    }
}

```

The Unique transformation uses a helper method nextInt which returns consecutive integers on each call. The basic transformation simply renames Var expressions. If, again, the expression language is extended with lambda constructs, the transformation needs to be updated as well to make the variable in the binding position of lambda expression unique. The following code shows how this can be done in a modular fashion:

```

interface UniqueWithLambda<E> extends Unique<E>,
    LamAlgTransform<E> {
    default E Lam(String x, E e) {
        return lamAlg().Lam(x + nextInt(), e);
    }
}

```

Note that the transformation uses the lamAlg() algebra (from LamAlgTransform), to create lambda expressions.

Figure 16 and 17 give a high level overview of query and transformation extension using the examples for FreeVars and Unique, respectively. In the case of queries, the abstract m() method will be shared by both the FreeVars and FreeVarsWithLambda interfaces. On the other hand, transformations are based on multiple base algebras, for sets of data type constructors (e.g., expAlg() and lamAlg()).

Note finally that, in the current implementation of **Shy** transformations, it is assumed that the language signatures




Figure 17. Extension of the Unique transformation.

`ExpAlg` and `LamAlg` are completely independent. This is however, not an essential requirement. An alternative design could have `LamAlg` be a proper extension of `ExpAlg` (i.e. `LamAlg<E>` **extends** `ExpAlg<E>`). In that case, the generated `LamAlgTransform` would need to refine the return type of the `expAlg()` method.

9. Shy Implementation

A new section.

Shy is implemented using the standard Java annotations framework (`javax.annotation`) packaged in the Java Development Kit. All of the generic traversals discussed in the previous sections – generic queries, generalized generic queries, generic transformations, and contextual generic transformations – are automatically generated by **Shy** for an object algebra interface annotated with `@Algebra`. Furthermore, **Shy** includes the monoid interface as well as several useful instances of the monoid interface.

The big advantage of using standard Java annotations is that the code generation of the generic traversals can be done transparently: users do not need to use or install a tool to generate that code. As a result **Shy** is as simple to use as a conventional library. Moreover, **Shy** works not only in many Java IDEs (like Eclipse and IntelliJ), but also from command line compilation.

Specifically, when an Object Algebra interface is annotated with `@Algebra`, **Shy** gets all the information from the interface (i.e. the names and types in each method). Such information is used together with those generic templates presented before, to generate generic traversal code automatically. The generated code lives in packages “query” (queries and generalized queries), “transform” (transformations and contextual transformations), “util” (transformations as classes) and “combinator” (combinators for combining two Object Algebras).

With **Shy** framework, the code users need to write is much shorter, which can be observed from the number of methods. Suppose a data structure consists of m construc-

tors, and n operations are applied to the structure. A traditional OO implementation defines m classes for the constructors, and for operations, n methods are implemented in each class. With Object Algebras, n classes serve for operations, and each includes m methods for all cases. The same number of classes are required with **Shy** approach, nonetheless, only the monoid and the interesting part (usually one or several cases) are to be implemented for each operation, which makes the code much more concise.

10. Case Study

To illustrate the utility of **Shy** we have implemented a number of queries and transformations in the context of QL, a DSL for questionnaires which has been implemented using Object Algebras before [9]. QL is similar to MiniQL, except that it additionally features an if-then-else construct, computed questions (which will appear read only), and a richer expression language. For more information on the features of QL we refer to [7].

10.1 QL Queries and Transformations

The queries extract derived information from a QL program, such as the set of used variables, the data and control dependencies between questions, and the global type environment. The transformations include two transformations of language extensions to the base language. The first realizes a simple desugaring of “unless(c)...” to “if(not(c))...”. The second desugaring statically unfolds a constant bound loop construct (“repeat (i)...”) and renames variables occurring below it accordingly. Finally, we have implemented a simple rename variable operation, and a flattening normalizer which inlines the conditions of nested if-then constructs.

Table 1 shows the number of cases that had to be overridden to implement each particular operation. The top row shows the number of constructs for each syntactic category in QL (Exp, Stmt, and Form). As can be seen, none of the operations required implementing all cases. The last column shows the number of overridden cases as a percentage. For this set of queries and transformations, almost no expression cases needed to be overridden, except the “Var” case in collect variables, rename variable and desugar “repeat”⁴. The cases required for desugaring include the case of the language extension (e.g. Unless and Repeat, respectively). These cases are not counted in the total in the first row but are used to compute the percentage.

10.2 Chaining Transformations

A typical compiler consists of many transformations chained together in a pipeline. **Shy** transformations support this pattern by passing transformation algebras as the base algebra to the implementation of another transformation. For instance, the desugar unless transformation desugars the “un-

⁴Note, however, that the dependency extraction queries reuse the collect variables query on expressions.

Operation	Exp (18)	Stmt (5)	Form (1)	%
Collect variables	1			4%
Data dependencies		3	1	17%
Control dependencies		4	1	21%
Type environment		2		8%
Rename variable	1	2		13%
Inline conditions		4		17%
Desugar “unless”		1		4%
Desugar “repeat”	1	3		16%

Table 1. Number of overridden cases per query and transformation in the context of the QL implementation

less” statement to “if” statements in another algebra. The latter can represent yet another transformation.

In the context of QL, “unless” desugaring, condition inlining and variable renaming can be chained together as follows:

```
alg = new Desugar<>(new Inline<>(
    new Rename<>(Collections.singletonMap("x", "y"),
        new Format())));
```

The chained transformation `alg` first desugars “unless”, then inlines conditions, and finally renames `xs` to `ys`. The `Rename` transformation gets as base algebra an instance of `Format`, a pretty printer for QL.

The algebra `alg` can now be used to create questionnaires:

```
Function<IFormatWithPrecedence, IFormat> pp
= alg.form("myForm", Arrays.asList(
    alg.unless(alg.var("x"),
        alg.question("x", "X?", new TBoolean()))));
```

Since inlining is a contextual transformation, the result of constructing this simple questionnaire is a function object representing the “to be inlined” representation of the questionnaire after desugaring. The `IFormatWithPrecedence` and `IFormat` types are formatting operations, respectively representing expressions and statements; these types originate from the `Format` algebra passed to `Rename`.

Calling the function with a boolean expression representing `true` will trigger inlining of conditions and renaming. The result is then a formatting object (`IFormat`) which can be used to print out the transformed questionnaire:

```
form myForm { if (true && !y) y "X?" boolean }
```

As can be seen, the variable `x` has been renamed to `y`. The (renamed) condition `y` is now negated, because of the desugaring of “unless”. Finally, the result of inlining conditions can be observed from the conjunction in the `if` statement.

10.3 Shy Performance vs Vanilla ASTs

We compared the performance characteristics of the operations implemented using **Shy** with respect to vanilla implementations based on ordinary AST classes with ordinary methods representing the transformations and queries. The operations were executed on progressively larger QL programs (up to 140Kb). In the vanilla implementation, the




Figure 18. Performance comparison of control dependencies query.

program was parsed into an AST structure, and then the operation was invoked and measured. In the case of the **Shy** queries, constructing the “AST” corresponds to executing the query, so we measured that. For context-dependent transformations, however, building the “AST” corresponds to constructing the function to execute the transformation, hence we only measured invoking this function. The vanilla query implementations use the same monoid structures as in **Shy**.

The comparison of the control dependencies query is shown in Figure 18. The plot shows that the performance is quite comparable. On average, the **Shy** implementation of the query seems a little slower. This is probably caused by the extensive use of interfaces in the **Shy** framework, whereas the AST-based implementation only uses abstract and concrete classes. For transformations the performance difference is slightly more pronounced. Figure 19 shows the performance comparison of the inline conditions transformation. The greater difference can be explained by the fact that creating a new structure in a **Shy** transformation involves dynamically dispatched method calls instead of statically bound constructor calls.

11. Related Work

Structure-shy traversals have been an active research topic. Our approach to structure shy traversal is unique in that it supports separate compilation, modular type checking and data type extension. Furthermore, it can be applied in mainstream languages such as Java. While some approaches support some of these features, to the best of our knowledge, no approach supports all of them.

Adaptive Object-Oriented Programming (AOOP) AOOP is an extension of object-oriented programming aimed at increasing the flexibility and maintainability of programs [17].




Figure 19. Performance comparison of inline conditions transformation.

AOOP promotes the idea of structure-shyness to achieve those goals. In AOOP it is possible to select parts of a structure that should be visited. This is useful to do traversals on complex structures and focus only on the interesting parts of the structure relevant for computing the final output. The original approach to AOOP was based on a domain-specific language [17]. DJ is an implementation of AOOP in Java using reflection [21]. More recently DemeterF [5] improved previous approaches to AOOP by providing support for type-safe traversals, generics and data-generic function generation. **Shy** shares with AOOP the use of structure-shyness as a means to increase flexibility and adaptability of programs. Most AOOP approaches, however, are not type-safe. The exception is DemeterF where a custom type system was designed to ensure type-safety of generic functions. Unlike DemeterF, which is a separate language, **Shy** is a Java library. Moreover, compilation of DemeterF programs is implemented through static weaving, and thus appears to preclude separate compilation.

Strategic Programming Strategic programming is an approach to data structure traversal, which originated in term rewriting [3, 24, 25]. Computations are represented by simple, conditional rewrite rules, but the application of such rules is controlled separately using the concept of a strategy. Strategies can be primitive (e.g., “fail”) or composed using combinators (e.g., “try s else s' ”). The strategy concept has been ported to other paradigms. JJTRAVELER is an OO framework for strategic programming [27]. Lämmel et al. introduced typed strategy combinators in Haskell [14]. The relation between strategic programming and AOOP has been explored in [16].

The key tenet of strategic programming is separation of concerns: actual computation and traversal are specified separately. In **Shy**, the traversal of a data structure is also specified separately (in a super-interface), however, it is fixed

```
interface AccuTrafoExp<M,E> extends
  ExpAlg<Pair<M,E>> {
    Monoid<M> m(); ExpAlg<E> expAlg();
    default Pair<M,E> Var(String s) {
        return new Pair<>(m().empty(),
            expAlg().Var(s));
    }
    default Pair<M,E> Lit(int i) {
        return new Pair<>(m().empty(),
            expAlg().Lit(i));
    }
    default Pair<M,E> Add(Pair<M,E> e1, Pair<M,E> e2)
    {
        return new Pair<>(m().join(e1.fst(), e2.fst()),
            expAlg().Add(e1.snd(), e2.snd()));
    }
}
```

Figure 20. Combining query and transformation.

for specific styles of queries and transformations. For instance, both queries and transformations employ an innermost, bottom-up strategy.

The distinction between queries and transformations also originates from existing work in strategic programming. Lämmel et al. [14] discuss type unifying and type preserving traversals. Type unifying traversals correspond to queries, where all data type constructors are unified into a single monoid. Analogously, **Shy** transformations are type preserving in the sense that a transformation is an algebra which maps constructor calls to another algebra of the same type.

The ASF+SDF program transformation system distinguishes transforming and accumulating traversals, which correspond to our transformations and queries, respectively [24]. Furthermore, an *accumulating transforming* traversal combines both styles, by tupling the accumulated result and the transformed tree. This combination could easily be generated by **Shy** by having the boilerplate code construct the monoid and the transformed term in parallel (see Fig. 20).

Structure-Shy Traversals with Visitors A standard way to remove boilerplate in OOP is to use *default visitors* [10]. Default visitors can be used in similar ways to our generic traversals. Many programmers using the visitor pattern create such default visitor implementations to avoid boilerplate code. There are two important differences to our work. Firstly conventional visitors are not extensible in a type-safe way. Secondly, with **Shy** the code for generic traversals is automatically generated, whereas with default visitors such code usually has to be implemented by hand.

Visser [27] adapted the strategy combinators of Stratego [25, 26] to combinators that operate on object-oriented Visitors [8]. The resulting framework JJTRAVELER solves the problem of entangling traversal control within the accept methods, or in the Visitors themselves (which only allow static specialization). A challenge not addressed by JJTRAV-

ELER is type safety of traversal code: either the combinators need to be redefined for each data type, or client code needs to cast the generic objects of type `AnyVisitable` to the specific type. Even if specific combinators would be generated, however, the traversed data types would not be extensible.

Another approach to improve upon the standard Visitor pattern is presented in Palsberg et al. [22]. This work particularly addressed the fact that traditional Visitors operate on a fixed set of classes. As a result, the data type can not be extended without changing all existing Visitors as well. The proposed solution is a generic `Walkabout` class which accesses sub-components of arbitrary data structures using reflection. Unfortunately, the heavy use of reflection makes `Walkabouts` significantly slower than traditional Visitors. The authors state that the `Walkabout` class could be generated to improve performance, but note that the addition of a class could trigger regeneration. As a result the pattern does not support separate compilation. Our solution obtains the same kind of default behavior for traversal, without losing extensibility, type safety, or separate compilation.

Whereas the `Walkabout` provides generic navigation over an object structure, this navigation can be programmed explicitly using *guides* [4]. Guides insert one level of indirection between recursing on the children of a node in `visit` methods: the guide decides how to proceed the traversal. Since guides need to define how to proceed for each type that will be visited, they suffer from the same extensibility problem as ordinary Visitors. Generic guides, on the other hand, are dynamically typed and use reflection to call appropriate `visit` methods. The `Walkabout` can be formulated as such a generic guide.

Structure-Shy Traversals in Functional Programming In functional programming there has been a lot of research on type-safe structure-shy traversals. Lämmel and Peyton Jones’ “Scrap your Boilerplate” (SyB) [11–13] series introduced a practical design pattern for doing generic traversals in Haskell. The simple queries and transformations in **Shy** were partly inspired by SyB. However SyB and **Shy** use very different implementation techniques. SyB is implemented in Haskell and relies on a run-time type-safe cast mechanism. This approach allows SyB traversals to be encoded once-and-forall using a single higher-order function called `gfoldl`. In contrast, in **Shy** Java annotations are used to generate generic traversals for each structure.

A drawback of SyB traversals is that they are notoriously slow, partly due to the use of the run-time cast [1]. Another notable difference between SyB and **Shy** is with respect to extensibility. While **Shy** supports extensibility of both traversals and structures, the original SyB approach did not support any extensibility. Only in later work, Lämmel and Peyton Jones proposed an alternative design for SyB, based on type classes [29]. This design supports extensibility of traversals, but not of the traversed structures.

Closest to **Shy** is an approach proposed by Lämmel et al. [15] for dealing with the so-called “*large bananas*”. A large banana corresponds to the fold algebra of a complex structure. Object Algebras, which we use in our work, are an OO encoding of fold algebras [18, 19]. However Lämmel et al. work has not dealt with extensibility. Interestingly in their future work Lämmel et al. did mention that they would like “to cope with incomplete or extensible systems of datatypes”.

Language Extensions for Queries Inspired by XPath/X-Query there has been some work on adding support for similar types of queries on object-oriented structures. For example the work on *Cω* [2] extends C# with generalized member access, which allows simple XPath-like path expressions. Thus in *Cω* it is possible to express queries quite concisely. However, in contrast to **Shy** *Cω* is a language extension and it does not deal with transformations.

Object Algebras **Shy** traversals are based on Object Algebras [18]. The original motivation for object algebras was as a design pattern for OO programming that allowed improved extensibility and modularity of programs. Using object algebras it is possible to solve the well-known “Expression Problem” [28]. Later work [20, 23] has explored the use of *Object Algebra combinators*, and generalizations of object algebras to improve expressiveness and modularity. In particular it has been claimed that Object Algebras can be used to do *feature-oriented programming* [20], and to encode *attribute grammars* [23]. One domain where Object Algebras are especially useful is in the implementation of (extensible) languages. The QL language used in our case study is based on Gouseti et al. [9]. That work provides a realistic implementation of an extensible *domain-specific language* using Object Algebras. In contrast to our work, which focus on removing boilerplate code, previous work on Object Algebras was mostly motivated by improved programming support for extensibility and modularity. The combination of extensibility (inherited for free from Object Algebras) and structure-shy type-safe traversals adds a new dimension to our work that, as far as we know, has not been explored previously.

12. Conclusion

This paper showed how various types of default traversals for complex structures can be automatically provided by **Shy**. **Shy** traversals are written directly in Java and are type-safe, extensible and separately compilable. There has always been a tension between the correctness guarantees of static typing, and the flexibility of untyped/dynamically-typed approaches. **Shy** shows that even in type systems like Java’s, it is possible to get considerable flexibility and adaptability for the problem of boilerplate code in traversals of complex structures, without giving up modular static typing.

There are many avenues for future work. One area of research is to extend **Shy** traversals to support flexible traversal strategies, similarly to strategic programming [3, 24, 25].

Another line of work worth exploring is to adopt generalizations of object algebras [20] for added expressiveness of **Shy** traversals.

Year of ref[18] fixed.

Name of ref[10] fixed.

References

- [1] M. D. Adams and T. M. DuBuisson. Template your boilerplate: Using Template Haskell for efficient generic programming. In *Proceedings of the 2012 ACM SIGPLAN Haskell symposium*, Haskell '12, pages 13–24, 2012.
- [2] G. M. Bierman, E. Meijer, and W. Schulte. The essence of data access in comega. In *ECOOP*, 2005.
- [3] P. Borovanský, C. Kirchner, H. Kirchner, P.-E. Moreau, and M. Vittek. ELAN: A logical framework based on computational systems. *Electronic Notes in Theoretical Computer Science*, 4:35–50, 1996.
- [4] M. Bravenboer and E. Visser. Guiding visitors: Separating navigation from computation. Technical Report UU-CS-2001-42, Institute of Information and Computing Sciences, Utrecht University, 2001.
- [5] B. Chadwick and K. Lieberherr. Weaving generic programming and traversal performance. In *AOSD'10*, 2010.
- [6] N. de Bruijn. Lambda calculus notation with nameless dummies, a tool for automatic formula manipulation, with application to the Church-Rosser theorem. *Indagationes Mathematicae (Proceedings)*, 75(5):381–392, 1972.
- [7] S. Erdweg, T. van der Storm, M. Völter, M. Boersma, R. Bosman, W. R. Cook, A. Gerritsen, A. Hulshout, S. Kelly, A. Loh, et al. The state of the art in language workbenches. In *Software Language Engineering*. Springer, 2013.
- [8] E. Gamma, R. Helm, R. Johnson, and J. Vlissides. *Design Patterns: Elements of Reusable Object-oriented Software*. Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 1995. ISBN 0-201-63361-2.
- [9] M. Gouseti, C. Peters, and T. v. d. Storm. Extensible language implementation with Object Algebras (short paper). In *GPCE'14*, 2014.
- [10] M. E. N. III. Variations on the visitor pattern, 1996.
- [11] R. Lämmel and S. P. Jones. Scrap your boilerplate: A practical design pattern for generic programming. In *TLDI'03*, 2003.
- [12] R. Lämmel and S. P. Jones. Scrap more boilerplate: Reflection, zips, and generalised casts. In *ICFP '04*, 2004.
- [13] R. Lämmel and S. P. Jones. Scrap your boilerplate with class: Extensible generic functions. In *ICFP '05*, 2005.
- [14] R. Lämmel and J. Visser. Typed combinators for generic traversal. In *Practical Aspects of Declarative Languages*, pages 137–154. Springer, 2002.
- [15] R. Lämmel, J. Visser, and J. Kort. Dealing with large bananas. In J. Jeuring, editor, *Workshop on Generic Programming*, Ponte de Lima, July 2000. Technical Report UU-CS-2000-19, Universiteit Utrecht.
- [16] R. Lämmel, E. Visser, and J. Visser. Strategic programming meets adaptive programming. In *AOSD '03*, 2003.
- [17] K. J. Lieberherr. *Adaptive Object Oriented Software: The Demeter Method with Propagation Patterns*. PWS Publishing, 1996.
- [18] B. C. d. S. Oliveira and W. R. Cook. Extensibility for the masses, practical extensibility with object algebras. In *ECOOP '12*, 2012.
- [19] B. C. d. S. Oliveira, M. Wang, and J. Gibbons. The visitor pattern as a reusable, generic, type-safe component. In *OOP-SLA'08*, 2008.
- [20] B. C. d. S. Oliveira, T. van der Storm, A. Loh, and W. R. Cook. Feature-oriented programming with object algebras. In *ECOOP '13*, 2013.
- [21] D. Orleans and K. J. Lieberherr. DJ: Dynamic adaptive programming in Java. In *Reflection 2001*. Springer-Verlag, 2001.
- [22] J. Palsberg and C. B. Jay. The essence of the visitor pattern. In *COMPSAC'98*, 1998.
- [23] T. Rendel, J. I. Brachthäuser, and K. Ostermann. From object algebras to attribute grammars. In *OOPSLA '14*, 2014.
- [24] M. G. J. van den Brand, P. Klint, and J. J. Vinju. Term rewriting with traversal functions. *ACM Trans. Softw. Eng. Methodol.*, 12(2):152–190, Apr. 2003.
- [25] E. Visser and Z.-e.-A. Benaissa. A core language for rewriting. *Electronic Notes in Theoretical Computer Science*, 15: 422–441, 1998.
- [26] E. Visser, Z.-e.-A. Benaissa, and A. Tolmach. Building program optimizers with rewriting strategies. In *ICFP '98*, 1998.
- [27] J. Visser. Visitor combination and traversal control. In *OOP-SLA '01*, 2001.
- [28] P. Wadler. The Expression Problem. Email, Nov. 1998. Discussion on the Java Genericity mailing list.
- [29] P. Wadler and S. Blott. How to make ad-hoc polymorphism less ad hoc. In *POPL '89*, 1989.

A. Appendix

A.1 Complete Code

A.1.1 OO Approach for usedVars and rename

Below is the complete code for Fig. 3 (left). It implements usedVars and rename in the QL example, as an OO approach.

(Renaming updated.)

```
class Form {
  String name;
  List<Stmt> body;
  Form(String id, List<Stmt> body) {
    this.name = id;
    this.body = new ArrayList<Stmt>(body);
  }
  Set<String> usedVars() {
    Set<String> vars = new HashSet<>();
    body.forEach(s -> vars.addAll(s.usedVars()));
    return vars;
  }
  Form rename(String n1, String n2) {
    List<Stmt> ss = new ArrayList<>();
    for (Stmt s: body) ss.add(s.rename(n1, n2));
    return new Form(name, ss);
  }
}

abstract class Stmt {
  abstract Set<String> usedVars();
  abstract Stmt rename(String n1, String n2);
}

class If extends Stmt {
  Exp cond;
  Stmt then;
  If(Exp cond, Stmt then) {
    this.cond = cond;
    this.then = then;
  }
  Set<String> usedVars() {
    Set<String> vars =
      new HashSet<>(cond.usedVars());
    vars.addAll(then.usedVars());
    return vars;
  }
  If rename(String n1, String n2) {
    return new If(cond.rename(n1, n2),
      then.rename(n1, n2));
  }
}

class Question extends Stmt {
  String name, label, type;
  Question(String n, String l, String t) {
    this.name = n;
    this.label = l;
    this.type = t;
  }
  Set<String> usedVars() {
    return emptySet();
  }
  Question rename(String n1, String n2) {
    String newN = name.equals(n1) ? n2 : name;
    return new Question(newN, label, type);
  }
}
```

```
}
```

```
abstract class Exp {
  abstract Set<String> usedVars();
  abstract Exp rename(String n1, String n2);
}
```

```
class Lit extends Exp {
  int n;
  Lit(int n) {
    this.n = n;
  }
  Set<String> usedVars() {
    return emptySet();
  }
  Lit rename(String n1, String n2) {
    return new Lit(n);
  }
}
```

```
class Var extends Exp {
  String x;
  Var(String name) {
    this.x = name;
  }
  Set<String> usedVars() {
    return Collections.singleton(x);
  }
  Var rename(String n1, String n2) {
    String newN = x.equals(n1) ? n2 : x;
    return new Var(newN);
  }
}
```

```
class GEq extends Exp {
  Exp lhs, rhs;
  GEq(Exp lhs, Exp rhs) {
    this.lhs = lhs;
    this.rhs = rhs;
  }
  Set<String> usedVars() {
    Set<String> vars = new HashSet<>(lhs.usedVars());
    vars.addAll(rhs.usedVars());
    return vars;
  }
  GEq rename(String n1, String n2) {
    return new GEq(lhs.rename(n1, n2),
      rhs.rename(n1, n2));
  }
}
```

A.1.2 Rename implementing the QLA1g interface

The following code gives the implementation of Rename that implements QLA1g in Section 3.2. (Renaming updated.)

```
class Rename<E, S, F> implements QLA1g<E, S, F> {
  private QLA1g<E, S, F> alg;
  private String n1, n2;
  public Rename(QLA1g<E, S, F> alg, String n1,
    String n2) {
    this.alg = alg; this.n1 = n1; this.n2 = n2;
  }
  public F Form(String id, List<S> stmts) {
    return alg.Form(id, stmts);
  }
}
```

```

}
public S If(E cond, S then) {
    return alg.If(cond, then);
}
public S Question(String id, String lbl,
    String type) {
    String newN = id.equals(n1) ? n2 : id;
    return alg.Question(newN, lbl, type);
}
public E Lit(int x) { return alg.Lit(x); }
public E Var(String name) {
    String newN = name.equals(n1) ? n2 : name;
    return alg.Var(newN);
}
public E GEq(E lhs, E rhs) {
    return alg.GEq(lhs, rhs);
}
}

```

A.1.3 QALgQuery: generated code

The generated code for QALgQuery by Shy in Fig. 4.

```

public interface QALgQuery<R>
    extends QALg<R, R, R> {

    Monoid<R> m();

    default R Form(java.lang.String p0,
        java.util.List<R> p1) {
        R res = m().empty();
        res = m().join(res, m().fold(p1));
        return res;
    }

    default R GEq(R p0, R p1) {
        R res = m().empty();
        res = m().join(res, p0);
        res = m().join(res, p1);
        return res;
    }

    default R If(R p0, R p1) {
        R res = m().empty();
        res = m().join(res, p0);
        res = m().join(res, p1);
        return res;
    }

    default R Lit(int p0) {
        R res = m().empty();
        return res;
    }

    default R Question(java.lang.String p0,
        java.lang.String p1, java.lang.String p2) {
        R res = m().empty();
        return res;
    }

    default R Var(java.lang.String p0) {
        R res = m().empty();
        return res;
    }
}

```

A.1.4 QALgTransform and QALgTrans: generated code

The code for QALgTransform and its class representation QALgTrans for use, generated by Shy. See Fig. 5.

```

public interface QALgTransform<A0, A1, A2>
    extends QALg<A0, A1, A2> {

    QALg<A0, A1, A2> qLAlg();

    @Override
    default A2 Form(java.lang.String p0,
        java.util.List<A1> p1) {
        return qLAlg().Form(p0, p1);
    }

    @Override
    default A0 GEq(A0 p0, A0 p1) {
        return qLAlg().GEq(p0, p1);
    }

    @Override
    default A1 If(A0 p0, A1 p1) {
        return qLAlg().If(p0, p1);
    }

    @Override
    default A0 Lit(int p0) {
        return qLAlg().Lit(p0);
    }

    @Override
    default A1 Question(java.lang.String p0,
        java.lang.String p1, java.lang.String p2) {
        return qLAlg().Question(p0, p1, p2);
    }

    @Override
    default A0 Var(java.lang.String p0) {
        return qLAlg().Var(p0);
    }
}

public class QALgTrans<A0, A1, A2>
    implements QALgTransform<A0, A1, A2> {

    private QALg<A0, A1, A2> alg;

    public QALgTrans(QALg<A0, A1, A2> alg) {
        this.alg = alg;
    }

    public QALg<A0, A1, A2> qLAlg() {return alg;}
}

```

A.1.5 G_ExpAlgQuery: generated code

The generated code for G_ExpAlgQuery by Shy in Fig. 10.

```

public interface G_ExpAlgQuery<A0>
    extends ExpAlg<A0> {

    Monoid<A0> mExp();

    @Override

```

```

default A0 Add(A0 p0, A0 p1) {
    A0 res = mExp().empty();
    res = mExp().join(res, p0);
    res = mExp().join(res, p1);
    return res;
}

@Override
default A0 Lit(int p0) {
    A0 res = mExp().empty();
    return res;
}

@Override
default A0 Var(java.lang.String p0) {
    A0 res = mExp().empty();
    return res;
}
}

```

A.1.6 G_ExpAlgTransform and G_LamAlgTransform: generated code

Below is the generated code for G_ExpAlgTransform and G_LamAlgTransform by Shy in Fig. 15.

```

public interface G_ExpAlgTransform<A, B0>
    extends ExpAlg<Function<A, B0>> {

    ExpAlg<B0> expAlg();

    default <B> List<B> substListExpAlg(List<Function<
        A, B>> list, A acc) {
        List<B> res = new ArrayList<B>();
        for (Function<A, B> i : list)
            res.add(i.apply(acc));
        return res;
    }

    @Override
    default Function<A, B0> Add(Function<A, B0> p0,
        Function<A, B0> p1) {
        return acc -> expAlg().Add(p0.apply(acc),
            p1.apply(acc));
    }

    @Override
    default Function<A, B0> Lit(int p0) {
        return acc -> expAlg().Lit(p0);
    }

    @Override
    default Function<A, B0> Var(java.lang.String p0) {
        return acc -> expAlg().Var(p0);
    }
}

public interface G_LamAlgTransform<A, B0>
    extends LamAlg<Function<A, B0>> {

    LamAlg<B0> lamAlg();

    default <B> List<B> substListLamAlg(List<Function<
        A, B>> list, A acc) {

```

```

        List<B> res = new ArrayList<B>();
        for (Function<A, B> i : list)
            res.add(i.apply(acc));
        return res;
    }

    @Override
    default Function<A, B0> Apply(Function<A, B0> p0,
        Function<A, B0> p1) {
        return acc -> lamAlg().Apply(p0.apply(acc),
            p1.apply(acc));
    }

    @Override
    default Function<A, B0> Lam(java.lang.String p0,
        Function<A, B0> p1) {
        return acc -> lamAlg().Lam(p0, p1.apply(acc));
    }
}

```

A.1.7 Util.cons: an auxiliary method

The auxiliary method Util.cons is implemented as follows, for the De Bruijn example in Fig. 15.

```

public class Util {
    public static <X> List<X> cons(X x, List<X> l) {
        l = new ArrayList<>(l);
        l.add(0, x);
        return l;
    }
}

```

A.1.8 Printer: a pretty printer for ExpAlg and LamAlg

The class Printer is implemented manually as a query algebra, for Section 7.1.

```

class Printer implements ExpAlg<String>,
    LamAlg<String> {
    public String Lam(String x, String e) {
        return "\"\" + x + \".\" + e;
    }
    public String Apply(String e1, String e2) {
        return "(" + e1 + " " + e2 + ")";
    }
    public String Var(String s) { return s; }
    public String Lit(int i) { return i + ""; }
    public String Add(String e1, String e2) {
        return "(" + e1 + " + " + e2 + ")";
    }
}

```